



US008812435B1

(12) **United States Patent**
Zhao

(10) **Patent No.:** **US 8,812,435 B1**
(45) **Date of Patent:** **Aug. 19, 2014**

(54) **LEARNING OBJECTS AND FACTS FROM DOCUMENTS**

(75) Inventor: **Shubin Zhao**, Jersey City, NJ (US)

(73) Assignee: **Google Inc.**, Mountain View, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 105 days.

5,787,413	A	7/1998	Kauffman et al.	707/2
5,793,966	A	8/1998	Amstein et al.	395/200.33
5,802,299	A	9/1998	Logan et al.	395/200.48
5,815,415	A	9/1998	Bentley et al.	
5,819,210	A	10/1998	Maxwell, III et al.	704/9
5,819,265	A	10/1998	Ravin et al.	
5,822,743	A	10/1998	Gupta et al.	
5,838,979	A	11/1998	Hart et al.	
5,909,689	A	6/1999	Van Ryzin	

(Continued)

FOREIGN PATENT DOCUMENTS

(21) Appl. No.: **11/941,382**

(22) Filed: **Nov. 16, 2007**

JP	5-174020	7/1993
JP	11-265400	9/1999

(Continued)

(51) **Int. Cl.**
G06F 7/00 (2006.01)

(52) **U.S. Cl.**
USPC **707/603; 706/45**

(58) **Field of Classification Search**
USPC **707/811, 603; 706/45**
See application file for complete search history.

OTHER PUBLICATIONS

Zhao et al., Corroborate and Learn Facts from the Web, Aug. 2007, AMC.*

(Continued)

(56) **References Cited**

U.S. PATENT DOCUMENTS

5,010,478	A	4/1991	Deran	364/200
5,133,075	A	7/1992	Risch	395/800
5,347,653	A	9/1994	Flynn et al.	
5,440,730	A	8/1995	Elmasri et al.	
5,475,819	A	12/1995	Miller et al.	395/200.03
5,519,608	A	5/1996	Kupiec	364/419.08
5,546,507	A	8/1996	Staub	
5,560,005	A	9/1996	Hoover et al.	395/600
5,574,898	A	11/1996	Leblang et al.	395/601
5,675,785	A	10/1997	Hall et al.	
5,680,622	A	10/1997	Even	
5,694,590	A	12/1997	Thuraisingham et al.	395/600
5,701,470	A	12/1997	Joy et al.	
5,717,911	A	2/1998	Madrid et al.	395/602
5,717,951	A	2/1998	Yabumoto	395/831
5,724,571	A	3/1998	Woods	
5,778,373	A	7/1998	Levy et al.	707/100
5,778,378	A	7/1998	Rubin	707/103

Primary Examiner — James Trujillo

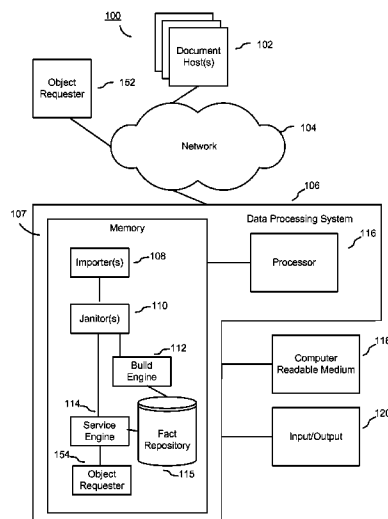
Assistant Examiner — Dawaune Conyers

(74) *Attorney, Agent, or Firm* — Morgan, Lewis & Bockius LLP

(57) **ABSTRACT**

A system, method, and computer program product for learning objects and facts from documents. A source object and a source document are selected and a title pattern and a contextual pattern are identified based on the source object and the source document. A set of documents matching the title pattern and the contextual pattern are selected. For each document in the selected set, a name and one or more facts are identified by applying the title pattern and the contextual pattern to the document. Objects are identified or created based on the identified names and associated with the identified facts.

20 Claims, 9 Drawing Sheets



(56)

References Cited

U.S. PATENT DOCUMENTS

5,920,859 A	7/1999	Li	707/5	6,886,005 B2	4/2005	Davis	707/2
5,943,670 A	8/1999	Prager	707/5	6,886,010 B2	4/2005	Kostoff	707/3
5,956,718 A	9/1999	Prasad et al.	707/10	6,901,403 B1	5/2005	Bata et al.	707/101
5,974,254 A	10/1999	Hsu		6,904,429 B2	6/2005	Sako et al.	707/3
5,987,460 A	11/1999	Niwa et al.	707/6	6,957,213 B1	10/2005	Yuret	707/4
6,006,221 A	12/1999	Liddy et al.	707/5	6,963,880 B1	11/2005	Pingte et al.	707/103
6,018,741 A	1/2000	Howland et al.		6,965,900 B2	11/2005	Srinivasa et al.	707/102
6,038,560 A	3/2000	Wical	707/5	7,003,506 B1	2/2006	Fisk et al.	707/1
6,044,366 A	3/2000	Graffe et al.	707/2	7,003,522 B1	2/2006	Reynar et al.	707/10
6,052,693 A	4/2000	Smith et al.	707/104	7,003,719 B1	2/2006	Rosenoff et al.	715/500.1
6,064,952 A	5/2000	Imanaka et al.	704/9	7,007,228 B1	2/2006	Carro	715/513
6,073,130 A	6/2000	Jacobson et al.	707/5	7,013,308 B1	3/2006	Tunstall-Pedoe	707/104.1
6,078,918 A	6/2000	Allen et al.		7,020,662 B2	3/2006	Boreham et al.	707/104.1
6,112,203 A	8/2000	Bharat et al.	707/5	7,043,521 B2	5/2006	Eitel	709/202
6,112,210 A	8/2000	Nori et al.		7,051,023 B2	5/2006	Kapur et al.	707/5
6,122,647 A	9/2000	Horowitz et al.	707/513	7,076,491 B2	7/2006	Tsao	707/100
6,134,555 A	10/2000	Chadha et al.	707/102	7,080,073 B1	7/2006	Jiang et al.	707/7
6,138,270 A	10/2000	Hsu		7,080,085 B1	7/2006	Choy et al.	707/101
6,182,063 B1	1/2001	Woods		7,100,082 B2	8/2006	Little et al.	
6,202,065 B1	3/2001	Wills	707/5	7,143,099 B2	11/2006	Lecheler-Moore et al.	707/101
6,212,526 B1	4/2001	Chaudhuri et al.	707/102	7,146,536 B2	12/2006	Bingham et al.	714/26
6,240,546 B1	5/2001	Lee et al.	717/4	7,158,980 B2	1/2007	Shen	707/100
6,263,328 B1	7/2001	Coden et al.	707/3	7,162,499 B2	1/2007	Lees et al.	707/203
6,263,358 B1	7/2001	Lee et al.		7,165,024 B2	1/2007	Glover et al.	704/9
6,266,805 B1	7/2001	Nwana et al.		7,174,504 B2	2/2007	Tsao	715/503
6,285,999 B1	9/2001	Page	707/5	7,181,471 B1	2/2007	Ibuki et al.	707/201
6,289,338 B1	9/2001	Stoffel et al.	707/3	7,194,380 B2	3/2007	Barrow et al.	702/181
6,311,194 B1	10/2001	Sheth et al.	707/505	7,197,449 B2	3/2007	Hu et al.	
6,314,555 B1	11/2001	Ndumu et al.		7,216,073 B2	5/2007	Lavi et al.	704/9
6,327,574 B1	12/2001	Kramer et al.	705/14	7,233,943 B2	6/2007	Modha et al.	
6,349,275 B1	2/2002	Schumacher et al.	704/8	7,260,573 B1	8/2007	Jeh et al.	
6,377,943 B1	4/2002	Jakobsson	707/2	7,263,565 B2	8/2007	Tawara et al.	
6,397,228 B1	5/2002	Lambert et al.		7,277,879 B2	10/2007	Varadarajan	707/1
6,438,543 B1	8/2002	Kazi et al.		7,302,646 B2	11/2007	Nomiyama et al.	715/764
6,470,330 B1	10/2002	Das et al.		7,305,380 B1	12/2007	Hoelzle et al.	707/3
6,473,898 B1	10/2002	Waugh et al.	717/168	7,325,160 B2	1/2008	Tsao	714/15
6,487,495 B1	11/2002	Gale et al.	701/209	7,363,312 B2	4/2008	Goldsack	707/102
6,502,102 B1	12/2002	Haswell et al.	707/102	7,376,895 B2	5/2008	Tsao	715/517
6,519,631 B1	2/2003	Rosenschein et al.	709/217	7,398,461 B1	7/2008	Broder et al.	
6,556,991 B1	4/2003	Borkovsky		7,409,381 B1	8/2008	Steel et al.	707/3
6,565,610 B1	5/2003	Wang et al.		7,412,078 B2	8/2008	Kim	
6,567,846 B1	5/2003	Garg et al.	709/218	7,418,736 B2	8/2008	Ghanea-Hercock	
6,567,936 B1	5/2003	Yang et al.	714/37	7,472,182 B1	12/2008	Young et al.	
6,572,661 B1	6/2003	Stern	715/501.1	7,483,829 B2	1/2009	Murakami et al.	
6,578,032 B1	6/2003	Chandrasekar et al.		7,493,308 B1	2/2009	Bair, Jr. et al.	707/3
6,584,464 B1	6/2003	Warthen	707/4	7,493,317 B2	2/2009	Geva	707/3
6,584,646 B2	7/2003	Fujita		7,587,387 B2	9/2009	Hogue	707/3
6,594,658 B2	7/2003	Woods		7,644,076 B1	1/2010	Ramesh et al.	
6,606,625 B1	8/2003	Muslea et al.	707/6	7,672,971 B2	3/2010	Betz et al.	707/103
6,606,659 B1	8/2003	Hegli et al.	709/225	7,685,201 B2	3/2010	Zeng et al.	
6,609,123 B1	8/2003	Cazemier et al.	707/4	7,698,303 B2	4/2010	Goodwin et al.	
6,636,742 B1	10/2003	Torkki et al.		7,716,225 B1	5/2010	Dean et al.	
6,643,641 B1	11/2003	Snyder	707/4	7,747,571 B2	6/2010	Boggs	
6,665,659 B1	12/2003	Logan	707/3	7,756,823 B2	7/2010	Young et al.	
6,665,666 B1	12/2003	Brown et al.	707/5	7,797,282 B1	9/2010	Kirshenbaum et al.	
6,665,837 B1	12/2003	Dean et al.	715/501.1	7,885,918 B2	2/2011	Statchuk	
6,675,159 B1	1/2004	Lin et al.		7,917,154 B2	3/2011	Fortescue et al.	
6,684,205 B1	1/2004	Modha et al.		7,953,720 B1	5/2011	Rohde et al.	
6,693,651 B2	2/2004	Biebesheimer et al.	345/837	8,024,281 B2	9/2011	Proctor et al.	
6,704,726 B1	3/2004	Amouroux	707/4	8,065,290 B2	11/2011	Hogue	
6,738,767 B1	5/2004	Chung et al.	707/7	8,108,501 B2	1/2012	Birnie et al.	
6,745,189 B2	6/2004	Schreiber		2001/0021935 A1	9/2001	Mills	707/513
6,754,873 B1	6/2004	Law et al.	715/501.1	2002/0022956 A1	2/2002	Ukrainczyk et al.	
6,763,496 B1	7/2004	Hennings et al.		2002/0038307 A1	3/2002	Obradovic et al.	707/102
6,799,176 B1	9/2004	Page	707/5	2002/0042707 A1	4/2002	Zhao et al.	704/9
6,804,667 B1	10/2004	Martin	707/6	2002/0065845 A1	5/2002	Naito et al.	707/500.1
6,820,081 B1	11/2004	Kawai et al.		2002/0073115 A1	6/2002	Davis	707/500.1
6,820,093 B2	11/2004	de la Hueraga		2002/0083039 A1	6/2002	Ferrari et al.	707/1
6,823,495 B1	11/2004	Vedula et al.	715/805	2002/0087567 A1	7/2002	Spiegler et al.	707/100
6,832,218 B1	12/2004	Emens et al.	707/3	2002/0107861 A1	8/2002	Clendinning et al.	707/101
6,845,354 B1	1/2005	Kuo et al.	704/9	2002/0147738 A1	10/2002	Reader	
6,850,896 B1	2/2005	Kelman et al.	705/10	2002/0169770 A1	11/2002	Kim et al.	707/5
6,868,411 B2	3/2005	Shanahan		2002/0174099 A1	11/2002	Raj et al.	707/1
6,873,982 B1	3/2005	Bates et al.	707/5	2002/0178448 A1	11/2002	Te Kieffe et al.	725/46
6,873,993 B2	3/2005	Charlesworth et al.		2002/0194172 A1	12/2002	Schreiber	707/4
				2003/0018652 A1	1/2003	Heckerman et al.	707/104.1
				2003/0058706 A1	3/2003	Okamoto et al.	365/200
				2003/0069880 A1	4/2003	Harrison et al.	707/3
				2003/0078902 A1	4/2003	Leong et al.	706/59

(56)

References Cited

U.S. PATENT DOCUMENTS

2003/0088607 A1	5/2003	Ruellan et al.	
2003/0097357 A1	5/2003	Ferrari et al.	707/3
2003/0120644 A1	6/2003	Shirota	
2003/0120675 A1	6/2003	Stauber et al.	707/100
2003/0126102 A1	7/2003	Borthwick	
2003/0126152 A1	7/2003	Rajak	707/101
2003/0149567 A1	8/2003	Schmitz et al.	704/270
2003/0149699 A1	8/2003	Tsao	707/100
2003/0154071 A1	8/2003	Shreve	704/9
2003/0167163 A1	9/2003	Glover et al.	
2003/0177110 A1	9/2003	Okamoto et al.	707/3
2003/0182310 A1	9/2003	Charnock et al.	707/104.1
2003/0195872 A1	10/2003	Senn	707/3
2003/0195877 A1	10/2003	Ford et al.	707/3
2003/0196052 A1	10/2003	Bolik et al.	
2003/0204481 A1	10/2003	Lau	
2003/0208354 A1	11/2003	Lin et al.	
2004/0003067 A1	1/2004	Ferrin	709/223
2004/0015481 A1	1/2004	Zinda	
2004/0024739 A1	2/2004	Copperman et al.	707/1
2004/0049503 A1	3/2004	Modha et al.	
2004/0059726 A1 *	3/2004	Hunter et al.	707/3
2004/0064447 A1	4/2004	Simske et al.	
2004/0069880 A1	4/2004	Samelson et al.	241/74
2004/0088292 A1	5/2004	Dettinger et al.	707/4
2004/0107125 A1	6/2004	Guheen et al.	705/7
2004/0122844 A1	6/2004	Malloy et al.	707/102
2004/0122846 A1	6/2004	Chess et al.	707/102
2004/0123240 A1	6/2004	Gerstl et al.	715/513
2004/0128624 A1	7/2004	Arellano et al.	715/530
2004/0143600 A1	7/2004	Musgrove et al.	
2004/0153456 A1	8/2004	Charnock et al.	707/10
2004/0167870 A1	8/2004	Wakefield et al.	707/1
2004/0167907 A1	8/2004	Wakefield et al.	
2004/0167911 A1	8/2004	Wakefield et al.	
2004/0177015 A1	9/2004	Galai et al.	705/35
2004/0177080 A1	9/2004	Doise et al.	707/100
2004/0199923 A1	10/2004	Russek	719/310
2004/0243552 A1	12/2004	Titemore et al.	707/3
2004/0243614 A1	12/2004	Boone et al.	
2004/0255237 A1	12/2004	Tong	715/501.1
2004/0260979 A1	12/2004	Kumai	
2004/0267700 A1	12/2004	Dumais et al.	
2004/0268237 A1	12/2004	Jones et al.	715/513
2005/0055365 A1	3/2005	Ramakrishnan et al.	
2005/0076012 A1	4/2005	Manber et al.	707/3
2005/0080613 A1	4/2005	Colledge et al.	
2005/0086211 A1	4/2005	Mayer	
2005/0086222 A1	4/2005	Wang et al.	707/5
2005/0086251 A1	4/2005	Hatscher et al.	
2005/0097150 A1	5/2005	McKeon et al.	
2005/0108630 A1 *	5/2005	Wasson et al.	715/513
2005/0125311 A1	6/2005	Chidiac et al.	705/28
2005/0149576 A1	7/2005	Marmaros et al.	707/200
2005/0149851 A1	7/2005	Mittal	
2005/0187923 A1	8/2005	Cipollone	707/3
2005/0188217 A1	8/2005	Ghanea-Hercock	
2005/0240615 A1	10/2005	Barsness et al.	707/102
2005/0256825 A1	11/2005	Dettinger et al.	707/1
2006/0036504 A1	2/2006	Allocca et al.	705/26
2006/0041597 A1	2/2006	Conrad et al.	
2006/0047691 A1	3/2006	Humphreys et al.	
2006/0047838 A1	3/2006	Chauhan	709/230
2006/0053171 A1	3/2006	Eldridge et al.	
2006/0053175 A1	3/2006	Gardner et al.	
2006/0064411 A1	3/2006	Gross et al.	707/3
2006/0074824 A1	4/2006	Li	706/20
2006/0074910 A1	4/2006	Yun et al.	707/7
2006/0085465 A1	4/2006	Nori et al.	707/101
2006/0112110 A1	5/2006	Maymir-Ducharme et al.	707/100
2006/0123046 A1	6/2006	Doise et al.	707/102
2006/0129843 A1	6/2006	Srinivasa et al.	
2006/0136585 A1	6/2006	Mayfield et al.	
2006/0143227 A1	6/2006	Helm et al.	

2006/0143603 A1	6/2006	Kalthoff et al.	
2006/0152755 A1	7/2006	Curtis et al.	358/1.15
2006/0167991 A1	7/2006	Heikes et al.	
2006/0224582 A1	10/2006	Hogue	707/6
2006/0238919 A1	10/2006	Bradley	
2006/0242180 A1	10/2006	Graf et al.	707/101
2006/0248045 A1	11/2006	Toledano et al.	
2006/0248456 A1	11/2006	Bender et al.	715/531
2006/0253418 A1	11/2006	Charnock et al.	707/1
2006/0259462 A1	11/2006	Timmons	707/3
2006/0277169 A1	12/2006	Lunt et al.	707/3
2006/0288268 A1	12/2006	Srinivasan et al.	
2006/0293879 A1 *	12/2006	Zhao et al.	704/9
2007/0005593 A1 *	1/2007	Self et al.	707/6
2007/0005639 A1	1/2007	Gaussier et al.	707/103
2007/0016890 A1	1/2007	Brunner et al.	717/107
2007/0038610 A1	2/2007	Omoigui	707/3
2007/0043708 A1	2/2007	Tunstall-Pedoe	
2007/0055656 A1	3/2007	Tunstall-Pedoe	707/3
2007/0073768 A1	3/2007	Goradia	707/104.1
2007/0094246 A1	4/2007	Dill et al.	707/4
2007/0100814 A1	5/2007	Lee et al.	
2007/0130123 A1	6/2007	Majumder	707/3
2007/0143282 A1	6/2007	Betz et al.	707/5
2007/0143317 A1 *	6/2007	Hogue et al.	707/100
2007/0150800 A1 *	6/2007	Betz et al.	715/500
2007/0198451 A1 *	8/2007	Kehlenbeck et al.	706/50
2007/0198480 A1	8/2007	Hogue et al.	707/3
2007/0198481 A1	8/2007	Hogue et al.	707/3
2007/0198503 A1	8/2007	Hogue et al.	707/5
2007/0198577 A1 *	8/2007	Betz et al.	707/103 R
2007/0198598 A1	8/2007	Betz et al.	707/201
2007/0198600 A1	8/2007	Betz	707/201
2007/0203867 A1	8/2007	Hogue et al.	706/48
2007/0208773 A1	9/2007	Tsao	707/103
2007/0271268 A1	11/2007	Fontoura et al.	707/6
2008/0071739 A1	3/2008	Kumar et al.	
2008/0104019 A1	5/2008	Nath	
2009/0006359 A1	1/2009	Liao	
2009/0119255 A1	5/2009	Frank et al.	707/3

FOREIGN PATENT DOCUMENTS

JP	2002-157276	5/2002	
JP	2002-540506	11/2002	
JP	2003-281173	10/2003	
WO	WO 01/27713	4/2001	
WO	WO 2004/114163	12/2004	G06F 17/30
WO	WO 2006/104951	10/2006	G06F 17/30

OTHER PUBLICATIONS

Agichtein, E., et al., "Snowball Extracting Relations from Large Plain-Text Collections," Columbia Univ. Computer Science Dept. Technical Report CUCS-033-99, Dec. 1999, pp. 1-13.

Brill, E., et al., "An Analysis of the AskMSR Question-Answering System," Proceedings of the Conference of Empirical Methods in Natural Language Processing (EMNLP), Jul. 2002, pp. 257-264.

Brin, S., Extracting Patterns and Relations from the World Wide Web, 1999, 12 pages.

Brin, S., et al., "The Anatomy of a Large-Scale Hypertextual Web Search Engine," 7th Int'l World Wide Web Conference, Brisbane, Australia, Apr. 14-18, 1998, pp. 1-26.

Bunescu, R., et al., "Using Encyclopedia Knowledge for Named Entity Disambiguation," Department of Computer Sciences, University of Texas, retrieved from internet Dec. 28, 2006, 8 pages.

Chang, C., et al., "IEPAD: Information Extraction Based on Pattern Discovery," WWW 10 '01, ACM, May 1-5, 2001, pp. 681-688.

Chu-Carroll, J., et al., "A Multi-Strategy and Multi-Source Approach to Question Answering," 2006, 8 pages.

Craswell, N., et al., "Effective Site Finding using Link Anchor Information," SIGIR '01, Sep. 9-12, 2001, pp. 250-257.

Dean, J., et al., "MapReduce: Simplified Data Processing on Large Clusters," OSDI, 2004, pp. 1-13.

Dong, X., et al., "Reference Reconciliation in Complex Information Spaces," SIGACM-SIGMOD, 2005, 12 pages.

(56)

References Cited

OTHER PUBLICATIONS

- Downey, D., et al., "Learning Text Patterns for Web Information Extraction and Assessment," American Association for Artificial Intelligence, 2002, 6 pages.
- Etzioni, O., et al., "Web-scale Information Extraction in KnowItAll (Preliminary Results)," WWW2004, ACM May 17-20, 2004, 11 pages.
- Freitag, D., et al., "Boosted Wrapper Induction," American Association for Artificial Intelligence, 2000, 7 pages.
- Gao, X., et al., "Learning Information Extraction Patterns from Tabular Web Pages Without Manual Labelling," Proceedings of IEEE/WIC Int'l Conf. on Web Intelligence (WI'03), Oct. 13-17, 2003, pp. 495-498.
- Gray, R.M., "Entropy and Information Theory," Springer-Verlag, New York, NY, 1990, pp. 17-46.
- Guha, R., "Object Co-Identification on the Semantic Web," WWW2004, ACM, May 17-22, 2004, 9 pages.
- Guha, R., et al., "Disambiguating People in Search," World Wide Web Conference, May 17-22, 2004, 9 pages.
- Haveliwala, T.H., "Topic-Sensitive PageRank," Proceeding of the 11th Int'l World Wide Web Conference, Honolulu, Hawaii, May 7-11, 2002, pp. 1-23.
- Hogue, A. W., Tree Pattern Inference and Matching for Wrapper Induction on the World Wide Web, Master of Engineering in Electrical Engineering and Computer Science, Massachusetts Institute of Technology, Jun. 2004, pp. 3-106.
- "Information Entropy," Wikipedia, the Free Encyclopedia, Retrieved on May 3, 2006, pp. 1-9.
- "Information Theory," Wikipedia, the Free Encyclopedia, Retrieved on May 3, 2006, pp. 1-12.
- International Search Report and Written Opinion for International Application No. PCT/US2007/61156, mailed Feb. 11, 2008, 7 pages.
- International Search Report and Written Opinion for International Application No. PCT/US2006/019807, mailed Dec. 18, 2006, 4 pages.
- International Search Report and Written Opinion for International Application No. PCT/US2006/07639, mailed Sep. 13, 2006, 6 pages.
- Jeh, G., et al., "Scaling Personalized Web Search," Proceedings of the 12th Int'l World Wide Web Conference, Budapest, Hungary, May 20-24, 2003, pp. 1-24.
- Ji, H., et al., "Re-Ranking Algorithms for Name Tagging," Workshop on Computationally Hard Problems and Joint Inference in Speech and Language Processing, Jun. 2006, 8 pages.
- Jones, R., et al., Bootstrapping for Text Learning Tasks, 1999, 12 pages.
- Kolodner, J., "Indexing and Retrieval Strategies for Natural Language Fact Retrieval," ACM Trans. Database Syst. 8.3., Sep. 1983, 434-464.
- Kosseim, L., et al., "Answer Formulation for Question-Answering," 11 pages, Oct. 1, 2007.
- Liu, B. et al., "Mining Data Records in Web Pages," Conference 2000, ACM, 2000, pp. 1-10.
- MacKay, D.J.C., "Information Theory, Inference and Learning Algorithms," Cambridge University Press, 2003, pp. 22-33, 138-140.
- Mann, G. et al., "Unsupervised Personal Name Disambiguation," Proceedings of the Seventy Conference on Natural Language Learning at HLT-NAACL, 2003, 8 pages.
- McCallum, A., et al., "Object Consolidation by Graph Partitioning with a Conditionally-Trained Distance Metric," SIGKDD 03, ACM, Aug. 24-27, 2003, 6 pages.
- Mihalcea, R., et al., PageRank on Semantic Networks, with Application to Word Sense Disambiguation, Proceedings of the 20th International Conference on Computational Linguistics, Aug. 23-27, 2004, 7 pages.
- Mihalcea, R., et al., "TextRank: Bringing Order into Texts," Proceedings of the Conference on Empirical Methods in Natural Language Processing, Jul. 2004, 8 pages.
- Page, L., et al., "The PageRank Citation Ranking: Bringing Order to the Web," Stanford Digital Libraries Working Paper, 1998, pp. 1-17.
- Pawson, D., "Sorting and Grouping," www.dpawson.co.uk/xsl/sect2/N6280.html, Feb. 7, 2004, pp. 1-19.
- Prager, J. et al., "IBM's Piquant in TREC2003," 2003, 10 pages.
- Prager, J., et al., "Question Answering Using Constraint Satisfaction: QA-by-Dossier-with-Constraints," 2004, 8 pages.
- Ramakrishnan, G., et al., "Is Question Answering an Acquired Skill?" WWW2004, ACM May 17, 2004, pp. 111-120.
- Richardson, M., et al., "Beyond Page Rank: Machine Learning for Static Ranking," International World Wide Web Conference Committee, May 23, 2006, 9 pages.
- Richardson, M., et al., "The Intelligent Surfer: Probabilistic Combination of Link and Content Information in PageRank," Advances in Neural Information Processing Systems, vol. 14, MIT Press, Cambridge, MA, 2002, 8 pages.
- Riolo, E., et al., "Learning Dictionaries for Information Extraction by Multi-Level Bootstrapping," American Association for Artificial Intelligence, 1999, 6 pages.
- Shannon, C.E., et al., "A Mathematical Theory of Communication," The Bell System Technical Journal, vol. 27, Jul., Oct. 1948, pp. 1-55.
- Sun Microsystems, "Attribute Names," <http://java.sun.com/products/jndi/tutorial/basics/directory/attnames.html>, Feb. 17, 2004, pp. 1-2.
- Wang, Y., et al., "C4-2: Combining Link and Contents in Clustering Web Search to Improve Information Interpretation," The University of Tokyo, 2002, pp. 1-9.
- Cover, T.M., et al., "Elements of Information Theory," Wiley-InterScience, New York, NY, 1991, pp. 12-23.
- Gigablast, Web/Directory, <http://www.gigablast.com/?c=dmoz3>, printed Aug. 24, 2010, 1 page.
- Gilster, P., "Get Fast Answers, Easily," The News Observer, May 14, 2003, 2 pages.
- Hsu, C. et al., "Finite-State Transducers for Semi-Structured Text Mining," IJCAI-99 Workshop on Text Mining: Foundations, Techniques and Applications, 1999, 12 pages.
- Ilyas, I. et al., "Rank-aware Query Optimization," SIGMOD 2004, Jun. 13-18, 2004, 12 pages.
- International Search Report and Written Opinion for International Application No. PCT/US2006/010965, mailed Jul. 5, 2006, 4 pages.
- Kosala, R., et al., "Web Mining Research: A Survey," SIGKDD Explorations, vol. 2, Issue 1, p. 1, Jul. 2000, 15 pages.
- Lin, J. et al., Question Answering from the Web Using Knowledge Annotation and Knowledge Mining Techniques, CIKM '03, Nov. 3-8, 2003, 8 pages.
- Nyberg, E. et al., "The JAVELIN Question-Answering System at TREC 2003: A Multi-Strategy Approach with Dynamic Planning," TREC 2003, 9 pages.
- Ogden, W. et al., "Improving Cross-Language Text Retrieval with Human Interactions," Proceedings of the 33rd Hawaii International Conference on System Sciences, IEEE 2000, 9 pages.
- Plaisant, C. et al., "Interface and Data Architecture for Query Preview in Networked Information Systems," ACM Transaction on Information Systems, vol. 17, Issue 3, Jul. 1999, 28 pages.
- Wirzenius, Lars, "C Preprocessor Trick for Implementing Similar Data Types," Jan. 17, 2000, pp. 1-9.
- Andritsos, Information-Theoretic Tools for Mining Database Structure from Large Data Sets, ACM SIGMOD, Jun. 13-18, 2004, 12 pgs.
- Chen, A Scheme for Inference Problems Using Rough Sets and Entropy, Lecture Notes in Computer Science, vol. 3642/2005, Regina, Canada, Aug. 31-Sep. 3, 2005, pp. 558-567.
- Dean, Using Design Recovery Techniques to Transform Legacy Systems, Software Maintenance, Nov. 7-9, 2001, Proceedings, IEEE International Conference, 10 pgs.
- Etzioni, Unsupervised Named-Entity Extraction from the Web: An Experimental Study, Dept. of Computer Science and Engineering, Univ. of Washington, Seattle, WA, Feb. 28, 2005, 42 pgs.
- Google, Canadian Patent Application 2610208, Office Action, Sep. 21, 2011, 3 pgs.
- Google, European Patent Application 06784449.8, Office Action, Mar. 26, 2012, 7 pgs.
- Google, Japanese Patent Application 2008-504204, Office Action, Oct. 12, 2011, 4 pgs.

(56)

References Cited**OTHER PUBLICATIONS**

Koeller, Approximate Matching of Textual Domain Attributes for Information Source Integration, IQIS '05 Proceedings of the 2nd International Workshop on Information Source Integration, Jun. 17, 2005, 10 pgs.

Merriam Webster Dictionary defines "normalize" as "to make conform to or reduce to a norm or standard", date: 1865, 2 pages.

Merriam Webster Dictionary defines "value" as "a numerical quantity that is assigned or is determined by calculation or measurement", date: 1300, 2 pages.

Microsoft Computer Dictionary defines "normalize" as "adjust number within specific range", May 1, 2002, 4 pages.

Microsoft Computer Dictionary Defines "quantity" as a "number", May 1, 2002, 4 pages.

Microsoft Computer Dictionary defines "value" as "a quantity", May 1, 2002, 4 pages.

Nadeau, Unsupervised named-Entity Recognition: Generating Gazetteers and Resolving Ambiguity, Inst. for Information Technology, National Research Council Canada, Gatineau and Ottawa, Canada, Aug. 1, 2006, 12 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Apr. 23, 2013, 21 pgs.

Betz, Office Action, U.S. Appl. No. 13/302,755, filed Mar. 25, 2012, 15 pgs.

Google Inc., Office Action, CA 2603085, Sep. 18, 2012, 2 pgs.

Hogue, Office Action, U.S. Appl. No. 13/549,361, filed Mar. 6, 2013, 13 pgs.

Hogue, Office Action, U.S. Appl. No. 13/603,354, filed Jan. 9, 2013, 5 pgs.

Laroco, Office Action, U.S. Appl. No. 13/364,244, filed Jan. 30, 2013, 8 pgs.

Shamsi, Office Action, U.S. Appl. No. 13/171,296, filed Apr. 3, 2013, 7 pgs.

Betz, Examiner's Answer, U.S. Appl. No. 11/097,688, filed Jul. 8, 2010, 18 pgs.

Betz, Examiner's Answer, U.S. Appl. No. 11/394,414, filed Jan. 24, 2011, 31 pgs.

Betz, Notice of Allowance, U.S. Appl. No. 11/142,740, filed Apr. 16, 2009, 7 pgs.

Betz, Notice of Allowance, U.S. Appl. No. 11/142,765, filed Jul. 1, 2010, 14 pgs.

Betz, Notice of Allowance, U.S. Appl. No. 11/341,069, filed Sep. 8, 2008, 6 pgs.

Betz, Notice of Allowance, U.S. Appl. No. 12/939,981, filed Aug. 11, 2011, 7 pgs.

Betz, Notice of Allowance, U.S. Appl. No. 12/939,981, filed Apr. 26, 2011, 11 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,740, filed Aug. 13, 2007, 12 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,740, filed May 17, 2007, 12 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,740, filed Jul. 23, 2008, 11 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,740, filed Dec. 26, 2007, 12 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,740, filed Jan. 27, 2009, 11 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,740, filed Apr. 30, 2008, 14 pgs.

Betz, Office Action, U.S. Appl. No. 11/097,688, filed Mar. 18, 2009, 13 pgs.

Betz, Office Action, U.S. Appl. No. 11/097,688, filed Oct. 29, 2009, 11 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,765, filed Jan. 8, 2010, 17 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,765, filed May 9, 2008, 20 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,765, filed Jan. 17, 2008, 16 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,765, filed Oct. 17, 2007, 14 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,765, filed Oct. 17, 2008, 17 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,765, filed Jun. 18, 2007, 13 pgs.

Betz, Office Action, U.S. Appl. No. 11/142,765, filed Apr. 28, 2009, 16 pgs.

Betz, Office Action, U.S. Appl. No. 11/341,069, filed Apr. 1, 2008, 8 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,414, filed Mar. 5, 2010, 24 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,414, filed Sep. 15, 2009, 16 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Apr. 1, 2008, 14 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Aug. 4, 2010, 19 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Feb. 8, 2011, 22 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Jul. 8, 2011, 13 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Apr. 11, 2012, 15 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Nov. 12, 2008, 11 pgs.

Betz, Office Action, U.S. Appl. No. 1/394,552, filed Jan. 13, 2010, 15 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Mar. 13, 2009, 12 pgs.

Betz, Office Action, U.S. Appl. No. 11/394,552, filed Sep. 24, 2012, 21 pgs.

Betz, Office Action, U.S. Appl. No. 12/939,981, filed Dec. 9, 2010, 12 pgs.

Hogue, Examiner's Answer, U.S. Appl. No. 11/142,748, filed Oct. 3, 2011, 23 pgs.

Hogue, Notice of Allowance, U.S. Appl. No. 11/907,689, filed Apr. 30, 2009, 8 pgs.

Hogue, Notice of Allowance, U.S. Appl. No. 11/356,837, filed Jan. 6, 2012, 12 pgs.

Hogue, Notice of Allowance, U.S. Appl. No. 11/356,837, filed Apr. 27, 2012, 7 pgs.

Hogue, Notice of Allowance, U.S. Appl. No. 12/546,578, filed Jan. 6, 2011, 8 pgs.

Hogue, Notice of Allowance, U.S. Appl. No. 12/546,578, filed Jul. 12, 2011, 10 pgs.

Hogue, Notice of Allowance, U.S. Appl. No. 13/206,457, filed Mar. 14, 2012, 9 pgs.

Hogue, Office Action, U.S. Appl. No. 11/907,689, filed Oct. 3, 2008, 13 pgs.

Hogue, Office Action, U.S. Appl. No. 11/097,689, Apr. 9, 2008, 11 pgs.

Hogue, Office Action, U.S. Appl. No. 11/097,689, filed Jun. 21, 2007, 9 pgs.

Hogue, Office Action, U.S. Appl. No. 11/097,689, Nov. 27, 2007, 10 pgs.

Hogue, Office Action, U.S. Appl. No. 11/142,748, filed Dec. 7, 2007, 13 pgs.

Hogue, Office Action, U.S. Appl. No. 11/142,748, filed Jul. 13, 2010, 12 pgs.

Hogue, Office Action, U.S. Appl. No. 11/142,748, Aug. 17, 2009, 14 pgs.

Hogue, Office Action, U.S. Appl. No. 11/142,748, Nov. 17, 2010, 14 pgs.

Hogue, Office Action, U.S. Appl. No. 11/142,748, filed May 18, 2007, 9 pgs.

Hogue, Office Action, U.S. Appl. No. 11/142,748, filed Jul. 22, 2008, 18 pgs.

Hogue, Office Action, U.S. Appl. No. 11/142,748, filed Aug. 23, 2007, 13 pgs.

Hogue, Office Action, U.S. Appl. No. 11/142,748, filed Jan. 27, 2009, 17 pgs.

(56)

References Cited

OTHER PUBLICATIONS

Hogue, Office Action, U.S. Appl. No. 11/356,837, filed Jun. 3, 2011, 18 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/356,837, filed Aug. 4, 2010, 20 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/356,837, filed Feb. 8, 2011, 14 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/356,837, filed May 11, 2009, 18 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/356,837, filed Feb. 19, 2010, 20 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/356,837, filed Mar. 21, 2008, 15 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/356,837, filed Oct. 27, 2009, 20 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/356,837, filed Sep. 30, 2008, 20 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/399,857, filed Mar. 1, 2012, 25 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/399,857, filed Mar. 3, 2011, 15 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/399,857, filed Jan. 5, 2009, 21 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/399,857, filed Jun. 8 2009, 14 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/399,857, filed Sep. 13, 2010, 13 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/399,857, filed Jun. 24, 2011, 14 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/399,857, filed Dec. 28, 2009, 11 pgs.
 Hogue, Office Action, U.S. Appl. No. 11/399,857, filed Mar. 31, 2008, 23 pgs.
 Hogue, Office Action, U.S. Appl. No. 12/546,578, filed Aug. 4, 2010, 10 pgs.
 Hogue, Office Action, U.S. Appl. No. 13/206,457, filed Oct. 28, 2011, 6 pgs.
 Hogue, Office Action, U.S. Appl. No. 13/549,361, filed Oct. 4, 2012, 18 pgs.
 Laroco, Notice of Allowance, U.S. Appl. No. 11/551,657, filed May 13, 2011, 8 pgs.
 Laroco, Notice of Allowance, U.S. Appl. No. 11/551,657, filed Sep. 28, 2011, 8 pgs.
 Laroco, Office Action, U.S. Appl. No. 11/551,657, filed Aug. 1, 2008, 15 pgs.
 Laroco, Office Action, U.S. Appl. No. 11/551,657, filed Aug. 13, 2009, 16 pgs.
 Laroco, Office Action, U.S. Appl. No. 11/551,657, filed Nov. 17, 2010, 20 pgs.
 Laroco, Office Action, U.S. Appl. No. 11/551,657, filed Feb. 24, 2010, 17 pgs.
 Laroco, Office Action, U.S. Appl. No. 11/551,657, filed Jan. 28, 2009, 17 pgs.
 Rohde, Notice of Allowance, U.S. Appl. No. 11/097,690, filed Dec. 23, 2010, 8 pgs.
 Rohde, Office Action, U.S. Appl. No. 11/097,690, filed May 1, 2008, 21 pgs.
 Rohde, Office Action, U.S. Appl. No. 11/097,690, filed Jun. 9, 2010, 11 pgs.
 Rohde, Office Action, U.S. Appl. No. 11/097,690, filed Oct. 15, 2008, 22 pgs.
 Rohde, Office Action, U.S. Appl. No. 11/097,690, filed Aug. 27, 2009, 13 pgs.

Rohde, Office Action, U.S. Appl. No. 11/097,690, filed Apr. 28, 2009, 9 pgs.
 Rohde, Office Action, U.S. Appl. No. 11/097,690, filed Sep. 28, 2007, 17 pgs.
 Shamsi, Notice of Allowance, U.S. Appl. No. 11/781,891, filed Oct. 25, 2010, 7 pgs.
 Shamsi, Notice of Allowance, U.S. Appl. No. 11/781,891, filed May 27, 2010, 6 pgs.
 Shamsi, Office Action, U.S. Appl. No. 11/781,891, filed Nov. 16, 2009, 10 pgs.
 Vespe, Notice of Allowance, U.S. Appl. No. 11/686,217, filed Aug. 27, 2012, 11 pgs.
 Vespe, Notice of Allowance, U.S. Appl. No. 11/745,605, filed Jun. 13, 2011, 9 pgs.
 Vespe, Notice of Allowance, U.S. Appl. No. 11/745,605, filed Sep. 22, 2011, 9 pgs.
 Vespe, Notice of Allowance, U.S. Appl. No. 11/745,605, filed Mar. 28, 2012, 10 pgs.
 Vespe, Office Action, U.S. Appl. No. 11/686,217, filed Sep. 10, 2010, 14 pgs.
 Vespe, Office Action, U.S. Appl. No. 11/686,217, filed Jan. 26, 2012, 12 pgs.
 Vespe, Office Action, U.S. Appl. No. 11/686,217, filed Mar. 26, 2010, 13 pgs.
 Vespe, Office Action, U.S. Appl. No. 11/745,605, filed Apr. 8, 2010, 15 pgs.
 Vespe, Office Action, U.S. Appl. No. 11/745,605, filed Jul. 30, 2009, 17 pgs.
 Zhao, Notice of Allowance, U.S. Appl. No. 11/394,610, filed May 11, 2009, 15 pgs.
 Zhao, Office Action, U.S. Appl. No. 11/142,853, filed Oct. 2, 2009, 10 pgs.
 Zhao, Office Action, U.S. Appl. No. 11/142,853, filed Sep. 5, 2008, 9 pgs.
 Zhao, Office Action, U.S. Appl. No. 11/142,853, filed Mar. 17, 2009, 9 pgs.
 Zhao, Office Action, U.S. Appl. No. 11/142,853, filed Jan. 25, 2008, 7 pgs.
 Zhao, Office Action, U.S. Appl. No. 11/394,610, filed Apr. 1, 2008, 18 pgs.
 Zhao, Office Action, U.S. Appl. No. 11/394,610, filed Nov. 13, 2008, 18 pgs.
 Betz, Final Office Action, U.S. Appl. No. 11/394,552, filed Oct. 21, 2013, 22 pgs.
 Betz, Notice of Allowance, U.S. Appl. No. 11/097,688, filed Nov. 19, 2013, 17 pgs.
 Betz, Notice of Allowance, U.S. Appl. No. 13/302,755, filed Aug. 28, 2013, 6 pgs.
 Hogue, Notice of Allowance, U.S. Appl. No. 13/549,361, filed Oct. 2, 2013, 9 pgs.
 Hogue, Notice of Allowance, U.S. Appl. No. 13/549,361, filed Jun. 26, 2013, 8 pgs.
 Hogue, Notice of Allowance, U.S. Appl. No. 13/603,354, filed Nov. 12, 2013, 9 pgs.
 Hogue, Notice of Allowance, U.S. Appl. No. 13/603,354, filed Jun. 26, 2013, 8 pgs.
 Laroco, Notice of Allowance, U.S. Appl. No. 13/364,244, filed Aug. 6, 2013, 6 pgs.
 Laroco, Office Action, U.S. Appl. No. 13/364,244, Dec. 19, 2013, 5 pgs.
 Shamsi, Final Office Action, U.S. Appl. No. 13/171,296, filed Nov. 4, 2013, 29 pgs.

* cited by examiner

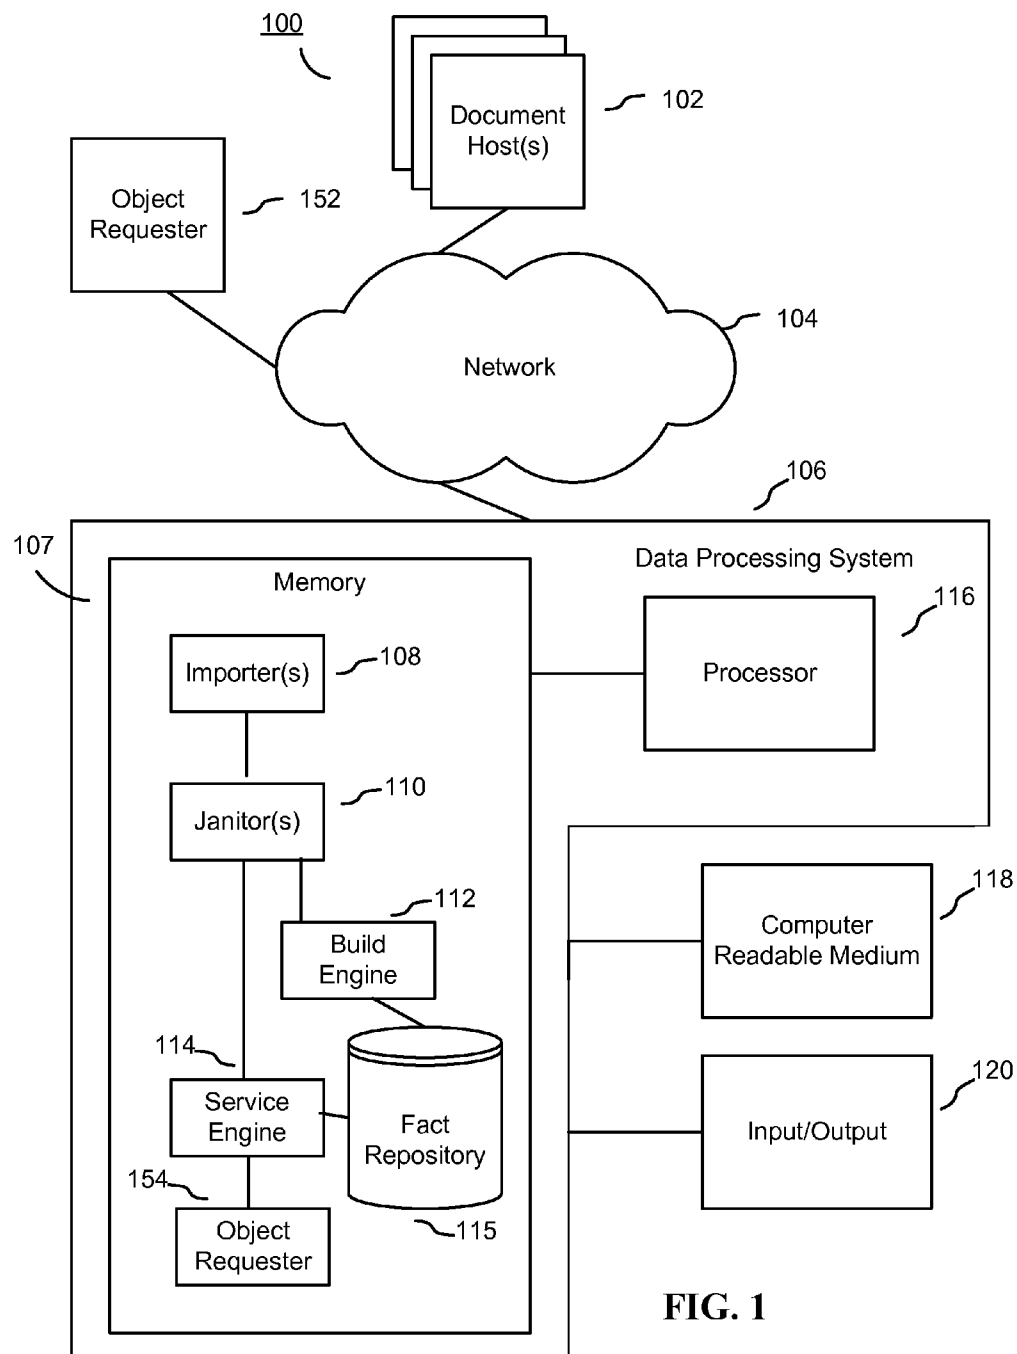


FIG. 1

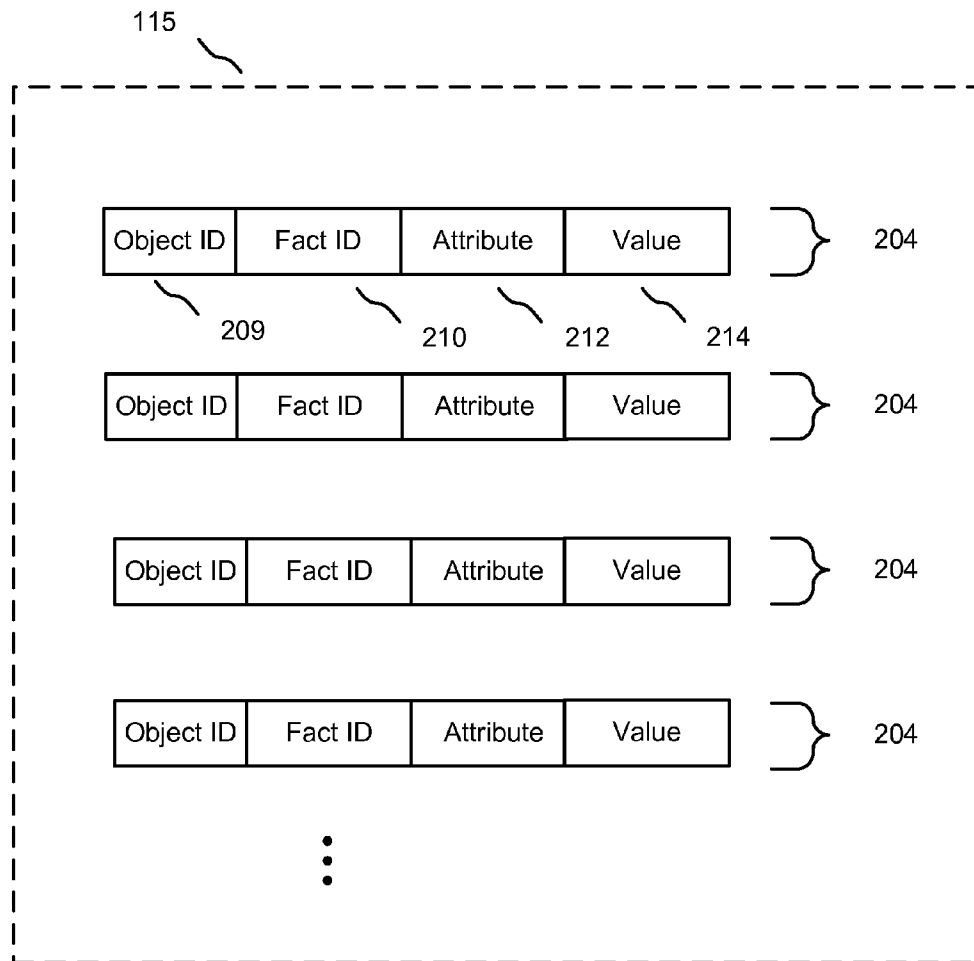


FIG. 2(a)
**Example Format of Facts in Repository (each
fact is associated with an object ID)**

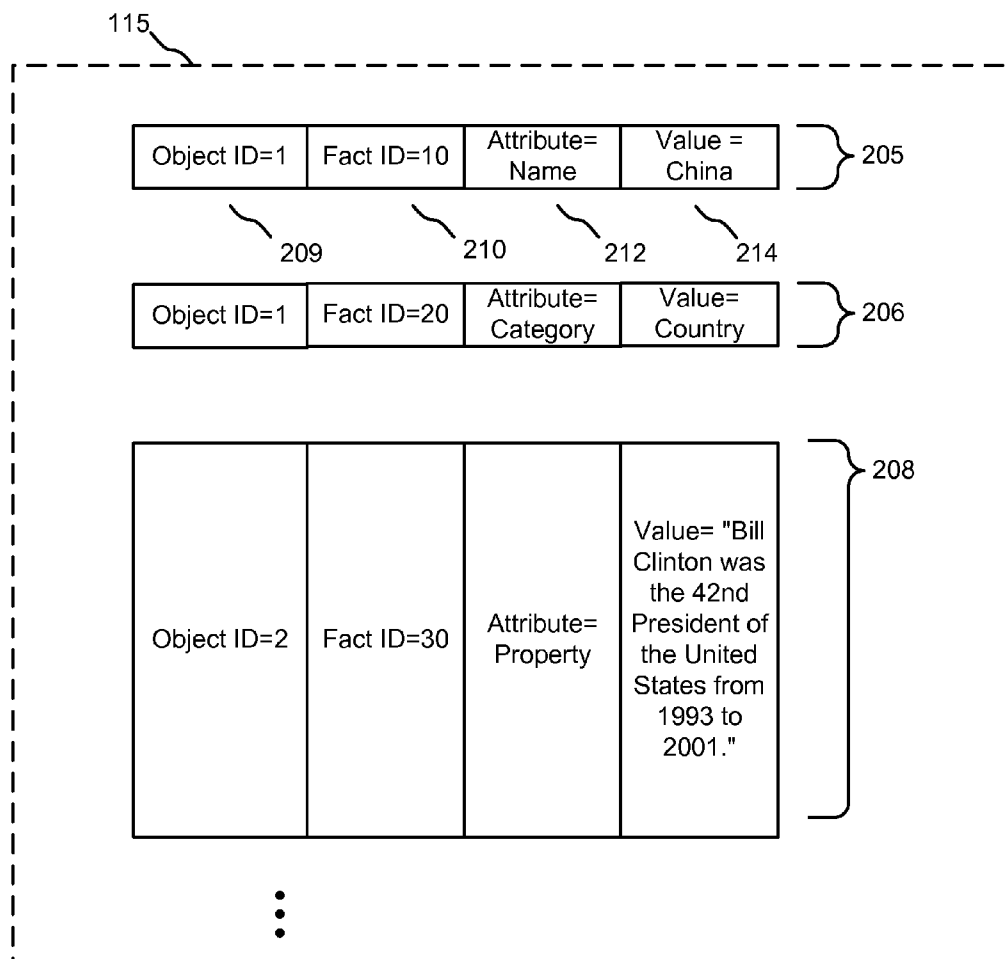
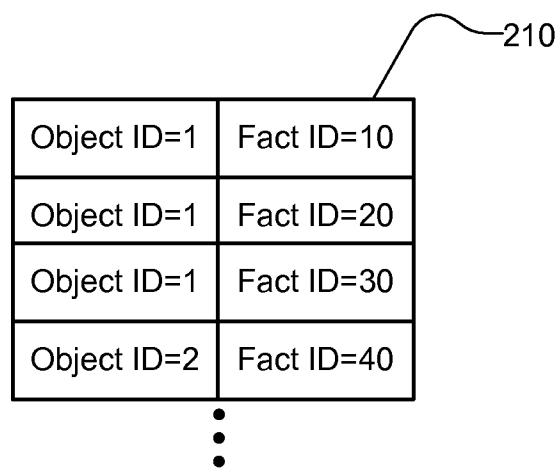


FIG. 2(b)
Example Facts in Repository (each fact is associated with an object ID)



Object ID=1	Fact ID=10
Object ID=1	Fact ID=20
Object ID=1	Fact ID=30
Object ID=2	Fact ID=40
⋮	

FIG. 2(c)
Example Object
Reference Table

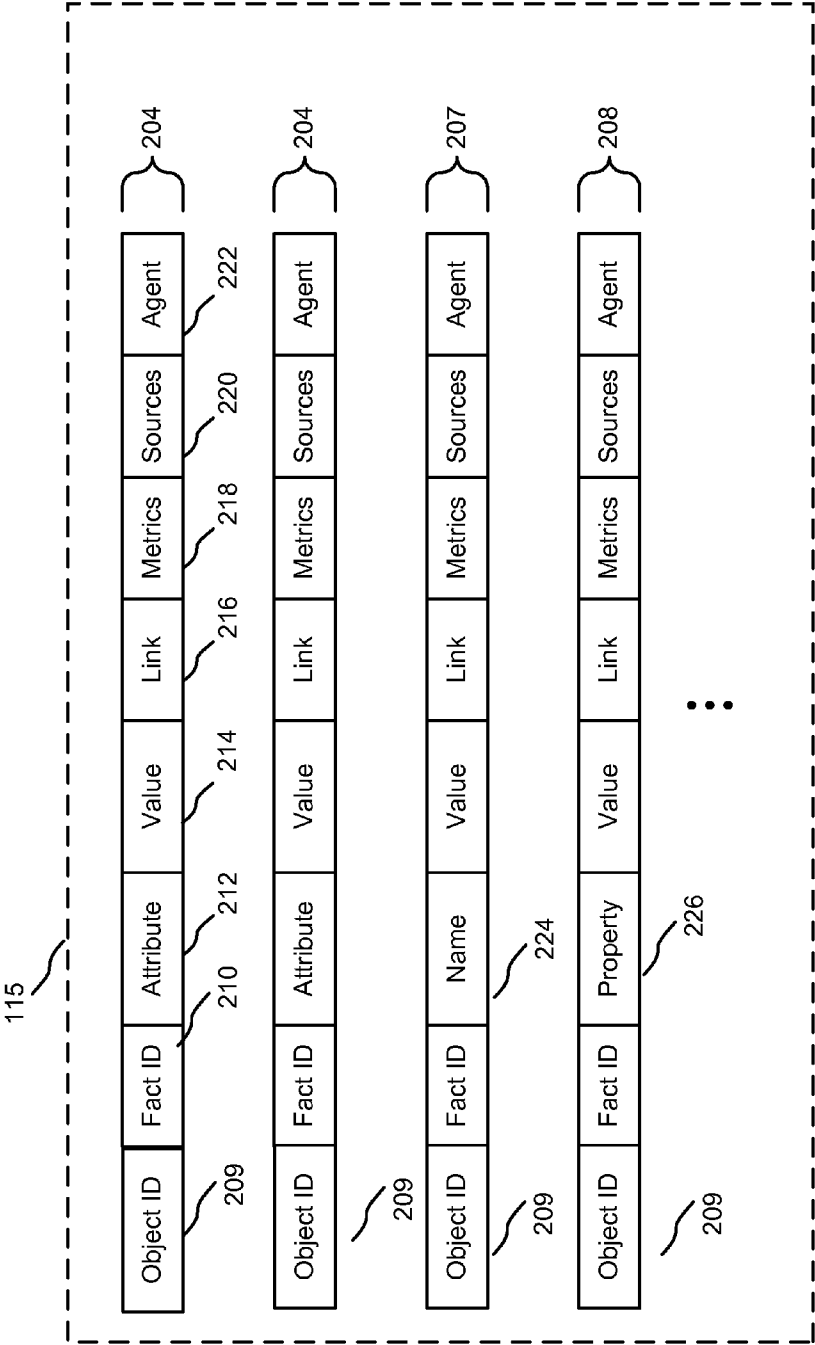


FIG. 2(d)
Example Format of Facts in
Repository (each fact is associated
with an object ID)

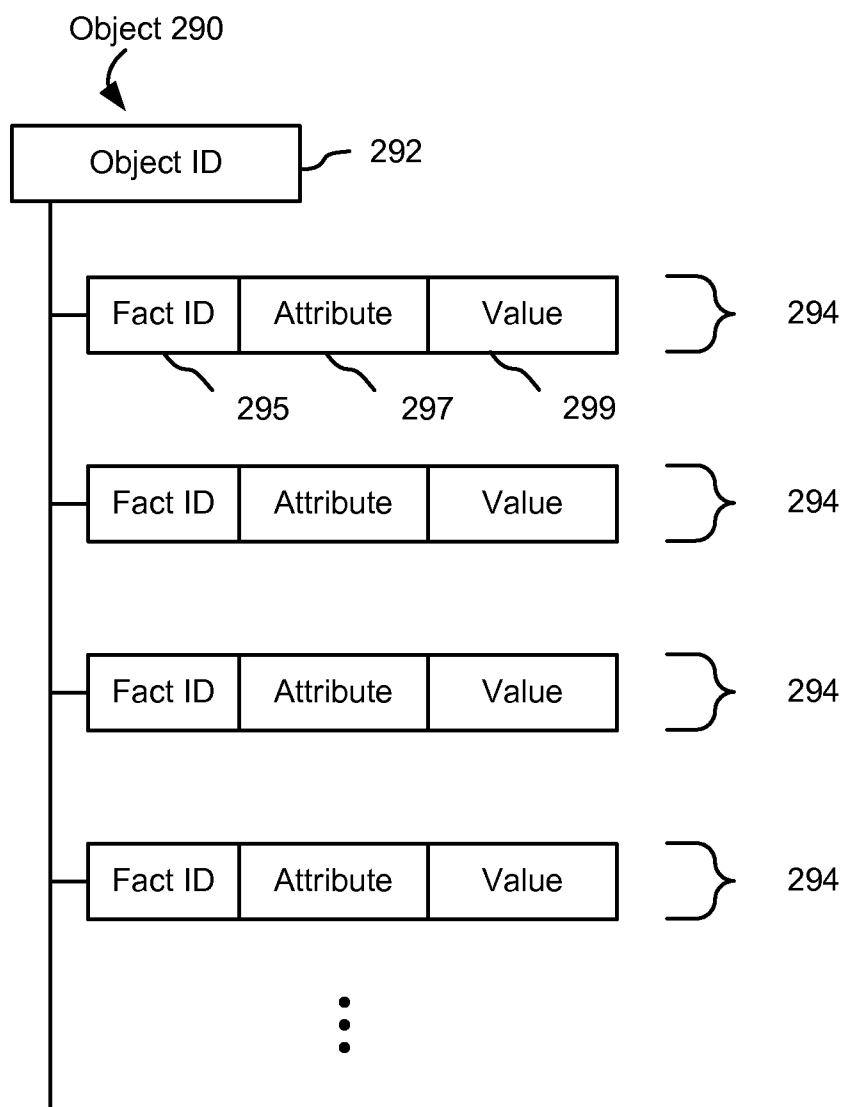
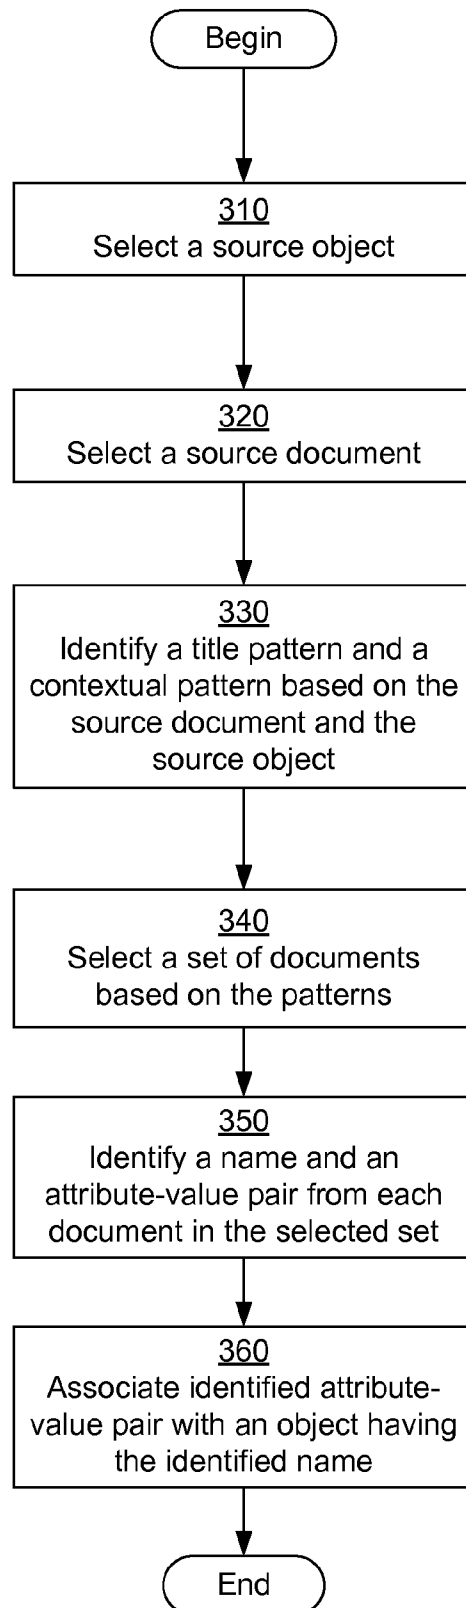
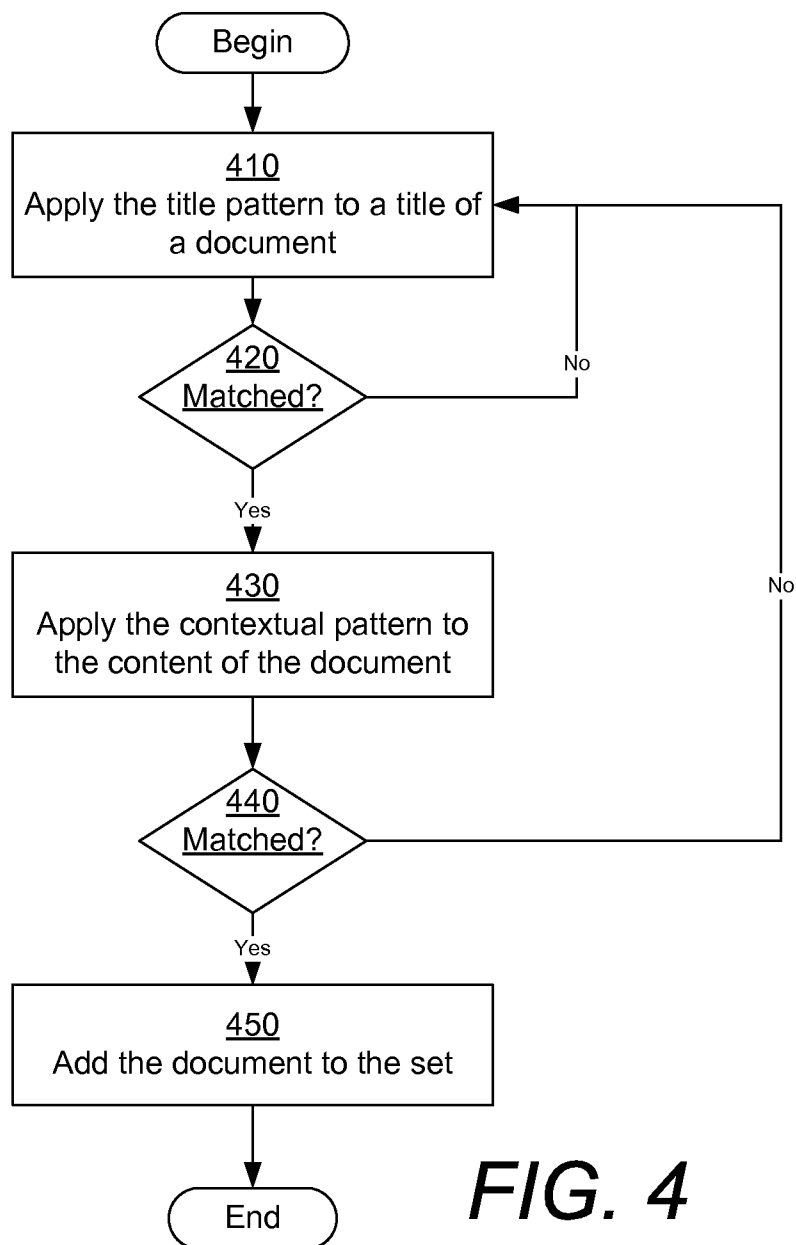
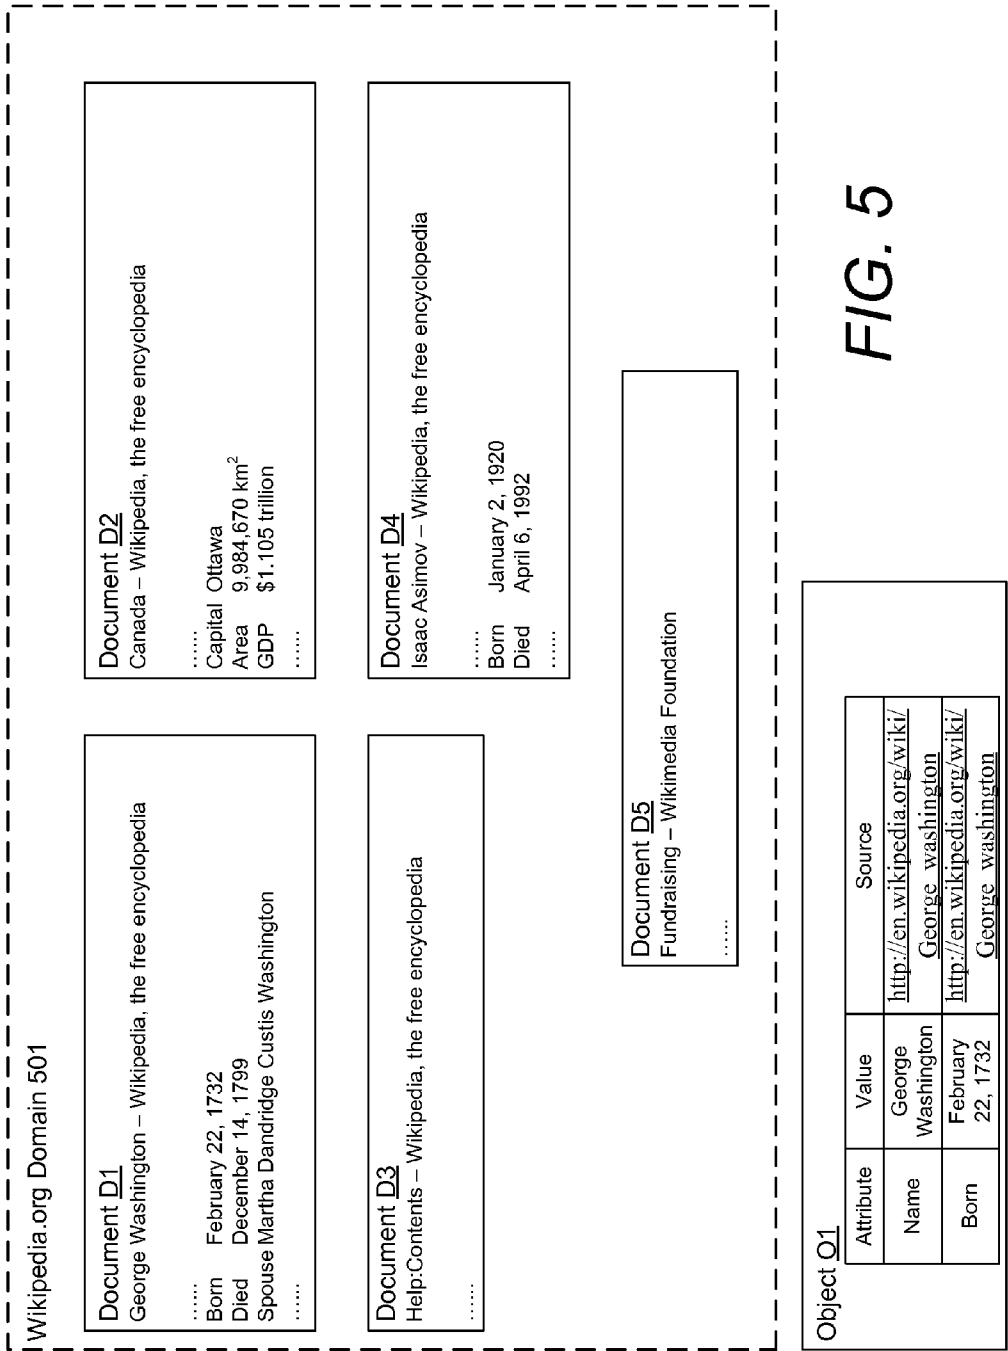


FIG. 2(e)
Example Objects

**FIG. 3**

**FIG. 4**



1

LEARNING OBJECTS AND FACTS FROM DOCUMENTS

CROSS-REFERENCE TO RELATED PATENT APPLICATIONS

This application is related to U.S. Utility patent application Ser. No. 11/394,610, entitled "Determining Document Subject by Using Title and Anchor Text of Related Documents," by Shubin Zhao, filed on Mar. 31, 2006, the disclosure of which is hereby incorporated by reference in its entirety.

TECHNICAL FIELD

The disclosed embodiments relate generally to fact databases. More particularly, the disclosed embodiments relate to learning objects and facts from documents.

BACKGROUND

As computers and networks gain popularity, web-based computer documents ("documents") become a vast source of factual information. Users may look to these documents to get answers to factual questions, such as "what is the capital of Poland" or "what is the birth date of George Washington." The factual information included in these documents may be extracted and stored in a fact database.

Documents are often generated based on a template. For example, titles of the documents in the wikipedia.org website often follow a pattern of "[SUBJECT]—Wikipedia, the free encyclopedia," where the section in square bracket is substituted with the subject of the 24207/11661/DOCS/1631371.6 page. These documents also often represent facts in a structured format. For example, documents in the wikipedia.org website frequently list facts in a table format.

Conventionally, objects (or entities) and related facts described in documents are identified and extracted (or learned) by human editors. This approach is insufficient for mass fact extraction because the vast volume of documents and the rapid increase in the number of available documents make it impractical for human editors to perform the task on any meaningful scale.

Based on the above, there is a need for a way to automatically identify objects and facts in documents.

SUMMARY

The above and other needs are met by systems, methods, and computer program products that learn objects and facts from documents. Embodiments of the methods comprise selecting a source object and a source document and identifying a title pattern and a contextual pattern based on the source object and the source document. A set of documents matching the title pattern and the contextual pattern are selected. For each document in the selected set, a name and one or more facts are identified by applying the title pattern and the contextual pattern to the document. Objects are identified or created based on the identified names and associated with the identified facts.

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 shows a system architecture adapted to support one embodiment.

FIGS. 2(a)-2(d) are block diagrams illustrating embodiments of a data structure for facts within a repository of FIG. 1.

2

FIG. 2(e) is a block diagram illustrating an embodiment of an alternative data structure for facts and objects.

FIG. 3 is a flowchart of an exemplary method for learning objects and facts from a plurality of documents in accordance with one embodiment.

FIG. 4 is a flowchart illustrating a process to identify a set of documents matching a title pattern and a contextual pattern in accordance with one embodiment.

FIG. 5 illustrates an example of learning objects and facts from a plurality of documents in accordance with one embodiment.

DESCRIPTION OF EMBODIMENTS

Embodiments are now described with reference to the figures where like reference numbers indicate identical or functionally similar elements.

System Architecture and Data Structure

FIG. 1 shows a system architecture 100 adapted to support one embodiment. FIG. 1 shows components used to add facts into, and retrieve facts from a repository 115. The system architecture 100 includes a network 104, through which any number of document hosts 102 communicate with a data processing system 106, along with any number of object requesters 152, 154.

Document hosts 102 store documents and provide access to documents. A document is comprised of any machine-readable data including any combination of text, graphics, multimedia content, etc. A document may be encoded in a markup language, such as Hypertext Markup Language (HTML), i.e., a web page, in an interpreted language (e.g., JavaScript) or in any other computer readable or executable format. A document can include one or more hyperlinks to other documents. A typical document will include one or more facts within its content. The facts describe entities, such as real-world or fictional people, places, or things.

A document stored in a document host 102 may be located and/or identified by a Uniform Resource Locator (URL), or Web address, or any other appropriate form of identification and/or location. A document host 102 is implemented by a computer system, and typically includes a server adapted to communicate over the network 104 via networking protocols (e.g., TCP/IP), as well as application and presentation protocols (e.g., HTTP, HTML, SOAP, D-HTML, JAVA®). The documents stored by a host 102 are typically held in a file directory, a database, or other data repository. A host 102 can be implemented in any computing device (e.g., from a PDA or personal computer, a workstation, mini-computer, or mainframe, to a cluster or grid of computers), as well as in any processor architecture or operating system.

FIG. 1 shows components used to manage facts in a fact repository 115. The data processing system 106 includes one or more importers 108, one or more janitors 110, a build engine 112, a service engine 114, and a fact repository 115 (also called simply a "repository"). Each of the foregoing is implemented, in one embodiment, as software modules (or programs) executed by the processor 116. Importers 108 operate to process documents received from the document hosts, read the data content of documents, and extract facts (as operationally and programmatically defined within the data processing system 106) from such documents. The importers 108 also determine the subject or subjects (i.e., the entity or entities) with which the facts are associated, and extract such facts into individual items of data, for storage in the repository 115. In one embodiment, there are different types of importers 108 for different types of documents, for example, dependent on the format or document type.

Janitors **110** operate to process facts extracted by the importer **108**. This processing can include but is not limited to, data cleansing, object merging, and fact induction. In one embodiment, there are a number of different janitors **110** that perform different types of data management operations on the facts. For example, one janitor **110** may traverse some set of facts in the repository **115** to find duplicate facts (that is, facts that convey the same factual information) and merge them. Another janitor **110** may also normalize facts into standard formats. Another janitor **110** may also remove unwanted facts from the repository **115**, such as facts related to pornographic content. Other types of janitors **110** may be implemented, depending on the types of data management functions desired, such as translation, compression, spelling or grammar correction, and the like.

Various janitors **110** act on facts to normalize attribute names and values, and delete duplicate and near-duplicate facts so an object does not have redundant information. For example, we might find on one page that Britney Spears' birthday is "Dec. 2, 1981" while on another page that her date of birth is "Dec. 2, 1981." Birthday and Date of Birth might both be rewritten as "Birthdate" by one janitor and then another janitor might notice that Dec. 2, 1981 and Dec. 2, 1981 are different forms of the same date. It would choose the preferred form, remove the other fact and combine the source lists for the two facts. As a result, one source page for this fact will contain an exact match of the fact while another source page will contain text that is considered synonymous with the fact.

The build engine **112** builds and manages the repository **115**. The service engine **114** is an interface for querying the repository **115**. The service engine **114**'s main function is to process queries, score matching objects, and return them to the caller but it is also used by the janitor **110**.

The repository **115** stores factual information about entities. The information is extracted from a plurality of documents that are located on document hosts **102**. A document from which a particular fact may be extracted is a source document (or "source") of that particular fact. In other words, a source of a fact includes that fact (or a synonymous fact) within its contents.

The repository **115** contains one or more facts. In one embodiment, the facts are logically organized into "objects," and each object contains a collection of facts associated with a single entity (i.e., real-world or fictional person, place, or thing). Each fact is associated with exactly one object. One implementation for this association includes in each fact an object ID that uniquely identifies the associated object. In this manner, any number of facts may be associated with an individual object, by including the object ID for that object in the facts. In one embodiment, objects themselves are not physically stored in the repository **115**, but rather are defined by the set or group of facts with the same associated object ID, as described below. Further details about facts in the repository **115** are described below, in relation to FIGS. 2(a)-2(d).

Some embodiments operate on the facts and/or objects in different orders than described above. For example, in one embodiment the importer **108** provides facts directly to the build engine **112** and/or repository **115**. The janitors **110**, in turn, operate on the facts and/or objects in the repository **115**. It should also be appreciated that in practice at least some of the components of the data processing system **106** will be distributed over multiple computers, communicating over a network. For example, the repository **115** may be deployed over multiple servers. As another example, the janitors **110** may be located on any number of different computers. For convenience of explanation, however, the components of the

data processing system **106** are discussed as though they were implemented on a single computer.

In another embodiment, some or all of document hosts **102** are located on the data processing system **106** instead of being coupled to the data processing system **106** by a network. For example, the importer **108** may import facts from a database that is a part of or associated with the data processing system **106**.

FIG. 1 also includes components to access the repository **115** on behalf of one or more object requesters **152**, **154**. Object requesters are entities that request objects from the repository **115**. Object requesters **152**, **154** may be understood as clients of the system **106**, and can be implemented in any computer device or architecture. As shown in FIG. 1, a first object requester **152** is located remotely from system **106**, while a second object requester **154** is located in the data processing system **106**. For example, in a computer system hosting a blog, the blog may include a reference to an object whose facts are in the repository **115**. An object requester **152**, such as a browser displaying the blog, will access data processing system **106** so that the information of the facts associated with the object can be displayed as part of the blog web page. As a second example, a janitor **110** or other entity considered to be part of data processing system **106** can function as an object requester **154**, requesting the facts of objects from the repository **115**.

FIG. 1 shows that the data processing system **106** includes a memory **107** and one or more processors **116**. The memory **107** includes the importers **108**, janitors **110**, build engine **112**, service engine **114**, and requester **154**, each of which is preferably implemented as instructions stored in memory **107** and executable by processor **116**. Memory **107** also includes the repository **115**. The repository **115** can be stored in a memory of one or more computer systems or in a type of memory such as a disk. FIG. 1 also includes a computer readable storage medium **118** containing, for example, at least one of importers **108**, janitors **110**, the build engine **112**, the service engine **114**, the requester **154**, and at least some portions of the repository **115**. FIG. 1 also includes one or more input/output devices **120** that allow data to be input and output to and from the data processing system **106**. It will be understood that embodiments of the data processing system **106** also include standard software components such as operating systems and the like and further include standard hardware components not shown in the figure for clarity of example.

FIG. 2(a) shows an example format of a data structure for facts within the repository **115**, according to some embodiments. As described above, the repository **115** includes facts **204** describing entities such as real-world and fictional people, places, and things. Each fact **204** includes a unique identifier for that fact, such as a fact ID **210**. Each fact **204** includes at least an attribute **212** and a value **214**. For example, a fact associated with the entity George Washington may include an attribute of "date of birth" and a value of "Feb. 22, 1732." In one embodiment, all facts are stored as alphanumeric characters since they are extracted from web pages. In another embodiment, facts also can store binary data values. Other embodiments, however, may store fact values as mixed types, or in encoded formats.

As described above, each fact is associated with an object ID **209** that identifies the object with which the fact is associated. Thus, each fact that describes the same entity (such as George Washington), will have the same object ID **209**. In one embodiment, the objects are logical concepts that exist as a collection of facts having the same object ID. In another embodiment, objects are stored as units of data in memory,

5

and include references (for example, pointers or IDs) to the facts associated with the object. The logical data structure of a fact can take various forms; in general, a fact is represented by a tuple that includes a fact ID, an attribute, a value, and an object ID. The storage implementation of a fact can be in any underlying physical data structure.

FIG. 2(b) shows an example of facts having respective fact IDs of 10, 20, and 30 in the repository 115. Facts 10 and 20 are associated with an object identified by object ID "1." Fact 10 has an attribute of "Name" and a value of "China." Fact 20 has an attribute of "Category" and a value of "Country." Thus, the object identified by object ID "1" has a name fact 205 with a value of "China" and a category fact 206 with a value of "Country." Fact 30 208 has an attribute of "Property" and a value of "Bill Clinton was the 42nd President of the United States from 1993 to 2001." Thus, the object identified by object ID "2" has a property fact with a fact ID of 30 and a value of "Bill Clinton was the 42nd President of the United States from 1993 to 2001." In the illustrated embodiment, each fact has one attribute and one value. The number of facts associated with an object is not limited; thus while only two facts are shown for the "China" object, in practice there may be dozens, even hundreds of facts associated with a given object. Also, the value fields of a fact need not be limited in size or content. For example, a fact about the economy of "China" with an attribute of "Economy" would have a value including several paragraphs of text, numbers, and perhaps even tables of figures. This content can be formatted, for example, in a markup language. For example, a fact having an attribute "original html" might have a value of the original html text taken from the source web page.

Also, while the illustration of FIG. 2(b) shows the explicit coding of object ID, fact ID, attribute, and value, in practice the content of the fact can be implicitly coded as well (e.g., the first field being the object ID, the second field being the fact ID, the third field being the attribute, and the fourth field being the value). Other fields include but are not limited to: the language used to state the fact (English, etc.), how important the fact is, the source of the fact, a confidence value for the fact, and so on.

FIG. 2(c) shows an example object reference table 210 that is used in some embodiments. Not all embodiments include an object reference table. The object reference table 210 functions to efficiently maintain the associations between object IDs and fact IDs. In the absence of an object reference table 210, it is also possible to find all facts for a given object ID by querying the repository 115 to find all facts with a particular object ID. While FIGS. 2(b) and 2(c) illustrate the object reference table 210 with explicit coding of object and fact IDs, the table also may contain just the ID values themselves in column or pair-wise arrangements.

FIG. 2(d) shows an example of a data structure for facts within the repository 115, according to some embodiments, showing an extended format of facts. In this example, the fields include an object reference link 216 to another object. The object reference link 216 can be an object ID of another object in the repository 115, or a reference to the location (e.g., table row) for the object in the object reference table 210. The object reference link 216 allows facts to have as values other objects. For example, for an object associated with the entity "United States," there may be a fact with the attribute of "president" and the value of "George W. Bush," with "George W. Bush" being an object having its own facts in the repository 115. In some embodiments, the value field 214 stores the name of the linked object and the link 216 stores the object identifier of the linked object. Thus, this "president" fact would include the value 214 of "George W.

6

Bush", and an object reference link 216 that contains the object ID for the "George W. Bush" object. In some other embodiments, facts 204 do not include a link field 216 because the value 214 of a fact 204 may store a link to another object.

Each fact 204 also may include one or more metrics 218. A metric provides an indication of the quality of the fact. In some embodiments, the metrics include a confidence level and an importance level. The confidence level indicates the likelihood that the fact is correct. The importance level indicates the relevance of the fact to the object, compared to other facts for the same object. The importance level may optionally be viewed as a measure of how vital a fact is to an understanding of the entity associated with the object.

Each fact 204 includes a list of one or more sources 220 that include the fact and from which the fact was extracted. Each source may be identified by a URL, or Web address, or any other appropriate form of identification and/or location, such as a unique document identifier.

The facts illustrated in FIG. 2(d) include an agent field 222 that identifies the importer 108 that extracted the fact. For example, the importer 108 may be a specialized importer that extracts facts from a specific source (e.g., the pages of a particular web site, or family of web sites) or type of source (e.g., web pages that present factual information in tabular form), or an importer 108 that extracts facts from free text in documents throughout the Web, and so forth.

Some embodiments include one or more specialized facts, such as a name fact 207 and a property fact 208. A name fact 207 is a fact that conveys a name for the entity associated with the object in which the fact is included. A name fact 207 includes an attribute 224 of "name" and a value, which is the name of the associated entity. For example, for an object associated with country Spain, a name fact would have the value "Spain." A name fact 207, being a special instance of a general fact 204, includes the same fields as any other fact 204; it has an attribute, a value, a fact ID, metrics, sources, etc. The attribute 224 of a name fact 207 indicates that the fact is a name fact, and the value is the actual name. The name may be a string of characters. An object may have one or more associated name facts, as many entities can have more than one name. For example, an object associated with Spain may have associated name facts conveying the country's common name "Spain" and the official name "Kingdom of Spain." As another example, an object associated with the U.S. Patent and Trademark Office may have associated name facts conveying the agency's acronyms "PTO" and "USPTO" as well as the official name "United States Patent and Trademark Office." If an object does have more than one associated name fact, one of the name facts may be designated as a primary name and other name facts may be designated as secondary names, either implicitly or explicitly. The name facts associated with an object are also called synonymous names of the object.

A property fact 208 is a fact that conveys a statement about the entity associated with the object. Property facts are generally used for summary information about an object. A property fact 208, being a special instance of a general fact 204, also includes the same fields (such as attribute, value, fact ID, etc.) as other facts 204. The attribute field 226 of a property fact 208 indicates that the fact is a property fact (e.g., attribute is "property") and the value is a string of text that conveys the statement of interest. For example, for the object associated with Bill Clinton, the value of a property fact may be the text string "Bill Clinton was the 42nd President of the United States from 1993 to 2001." Some objects may have one or more associated property facts while other objects may have

no associated property facts. It should be appreciated that the data structures shown in FIGS. 2(a)-2(d) and described above are merely exemplary. The data structure of the repository 115 may take on other forms. Other fields may be included in facts and some of the fields described above may be omitted. Additionally, each object may have additional special facts aside from name facts and property facts, such as facts conveying a type or category (for example, person, place, movie, actor, organization, etc.) for categorizing the entity associated with the object. In some embodiments, an object's name(s) and/or properties may be represented by special records that have a different format than the general fact records 204.

As described previously, a collection of facts is associated with an object ID of an object. An object may become a null or empty object when facts are disassociated from the object. A null object can arise in a number of different ways. One type of null object is an object that has had all of its facts (including name facts) removed, leaving no facts associated with its object ID. Another type of null object is an object that has all of its associated facts other than name facts removed, leaving only its name fact(s). Alternatively, the object may be a null object only if all of its associated name facts are removed. A null object represents an entity or concept for which the data processing system 106 has no factual information and, as far as the data processing system 106 is concerned, does not exist. In some embodiments, facts of a null object may be left in the repository 115, but have their object ID values cleared (or have their importance set to a negative value). However, the facts of the null object are treated as if they were removed from the repository 115. In some other embodiments, facts of null objects are physically removed from the repository 115.

FIG. 2(e) is a block diagram illustrating an alternate data structure 290 for facts and objects in accordance with embodiments of the invention. In this data structure, an object 290 contains an object ID 292 and references or points to facts 294. Each fact includes a fact ID 295, an attribute 297, and a value 299. In this embodiment, an object 290 actually exists in memory 107.

As described above, an object may explicitly exist in the repository 115, or it may exist merely as a collection of facts with a common object ID. Reference is made to particular objects for the purposes of illustration; one of skill in the art will recognize that the systems and methods described herein are applicable to a variety of implementations and that such references are not limiting. When reference is made to a fact being associated with an object, it should be understood that in at least one embodiment a fact is associated with an object by sharing a common object ID with other facts. For example, a fact could be associated with an object of a given type by sharing a common object ID at least with a type fact indicating the given type (or as another example, with a category fact indicating a particular category of object). Furthermore, in various embodiments, facts and objects can be stored in a variety of structures, such as fact and/or object repositories. When reference is made herein to the repository 115, it should be understood that various embodiments may store facts and/or objects in a variety of data structures.

Overview of Methodology

Referring now to FIG. 3, there is shown a flow diagram illustrating a method 300 for learning (or extracting) objects and facts from a plurality of documents in accordance with one embodiment. Other embodiments perform steps of the method 300 in different orders and/or perform different or additional steps than the ones shown in FIG. 3. The steps of

the process illustrated in FIG. 3 may be implemented in software, hardware, or a combination of hardware and software.

In one embodiment, the steps of the method 300 may be performed by one or more importers 108 as shown in FIG. 1, although one skilled in the art will recognize that the method could be performed by systems having different architectures as well. The importers 108 can perform multiple instances of the steps of the method 300 concurrently and/or perform steps in parallel.

The flowchart shown in FIG. 3 will now be described in detail. The process commences with a plurality of objects and a plurality of documents. The plurality of objects can be objects in the fact repository 115. The plurality of documents can be documents in the document hosts 102.

As described above in relation to FIGS. 2(a)-2(e), an object is associated with one or more facts. Each fact has an attribute and a value (hereinafter called an "attribute-value pair"). For example, an object representing Canada can have a fact with an attribute "name" and a value of the attribute "Canada". A fact with the attribute "name" is called a name fact, and its value is called an object name. The object representing Canada can also have a fact with an attribute-value pair of "Capital-Ottawa", wherein "Capital" is the attribute and "Ottawa" is the value.

As described above, a document may be encoded in a markup language, such as the Hypertext Markup Language (HTML), i.e., a web page, in an interpreted language (e.g., JavaScript) or in any other computer readable or executable format. A document can have one or more titles. One example is a web page encoded in HTML with content about the first President of the United States, George Washington.

The importer 108 selects 310 a source object from the plurality of objects. In one embodiment, the importer 108 gives priority to objects associated with more facts when selecting 310 the source object. In another embodiment, each of the plurality of objects has equal chance of being selected 310.

The importer 108 selects 320 a source document associated with a fact of the source object from the plurality of documents. As described above in relation to FIG. 2(d), each fact includes a list of sources 220 from which the fact was extracted. In one embodiment, the importer 108 selects 320 the source document from sources of a fact associated with the source object. In one embodiment, the importer 108 gives priority to documents from reputable websites (e.g., the Encyclopedia Britannica Online) when selecting 320 the source document.

Alternatively, the importer 108 searches for documents containing an object name of the source object and one or more of its facts (or attribute-value pairs) in the plurality of documents. The importer 108 then selects 320 the source document from the search result. For example, for an object with object name "China" and attribute-value pair "Capital-Beijing," the importer 108 can search for documents containing "China," "Capital," and "Beijing." The importer 108 may specify search criteria such as whether the object name matches document titles and whether the rest of the search terms match document contents. The search may be performed using a search engine, such as a Web search engine.

The importer 108 identifies 330 a title pattern and a contextual pattern based on the source document and the source object. The title pattern is intended to reflect a relationship between a title of the source document (hereinafter called "the document title") and the object name of the source object ("the source object name"). In one embodiment, the importer 108 determines whether the document title contains the

source object name as a substring. A substring is a contiguous sequence of characters taken from a string. If the document title does not contain the source object name as a substring, the importer **108** repeats the above process and selects **320** another document as the source document.

The importer **108** identifies **330** the title pattern by comparing the document title with the source object name. In one embodiment, a title pattern comprises a prefix section, a body section, and a suffix section. A title pattern can be expressed as “%[PREFIX SECTION]% && %[SUFFIX SECTION]”, where “&&” represents the source object name. For example, the title pattern of a document title “Summary for Microsoft CP—Yahoo! Finance” and an object name “Microsoft CP” is “%Summary for% && %—Yahoo! Finance%”. In this example, the prefix section is “Summary for” and the suffix section is “—Yahoo! Finance”. The body section of the document title is the same as the object name, “Microsoft CP.”

The contextual pattern is a structural pattern in which the source document displays (or presents) facts (or attribute-value pairs). For example, a document can display facts of an object in a table (the contextual pattern) such that the attribute-value pair for a fact appears in the same row. In one embodiment, the importer **108** identifies **330** the contextual pattern by identifying language markups (e.g., HTML markups (or tags)) that define the structural pattern. For example, a web page may contain a list of attributes and values with HTML markup as the following:

```
...
<title> China—Country Facts </title>
...
<table border="1">
<tr> <td> Population </td> <td> 1.3 billion </td> </tr>
<tr> <td> Capital </td> <td> Beijing </td> </tr>
<tr> <td> Largest City </td> <td> Shanghai </td> </tr>
</table>
...
```

The above web page contains facts about China such as that its capital is Beijing. The HTML markups indicate that the facts should be presented in a table format (e.g., by a web browser).

The importer **108** can identify **330** the contextual pattern in the source document by matching attribute-value pairs of the source object with content of the source document. The contextual pattern contains the HTML markup (or tag) sequence associated with the matching attribute-value. For example, for an attribute-value pair “Capital-Beijing,” the importer **108** may identify **330** the following contextual pattern in the above web page: “<tr> <td> \$ \$ </td> <td> £ £ </td> </tr>”. In this example, “\$ \$” indicates a text block for an attribute and “£ £” indicates a text block of a corresponding value. The importer **108** may identify **330** multiple contextual patterns in a source document, with each pattern corresponding with a fact of the source object.

The importer **108** selects **340** a set of documents from the plurality of documents based on the title pattern and the contextual pattern. One embodiment of the process for the importer **108** to select **340** the set of documents is further illustrated in the flowchart shown in FIG. 4. For each of the plurality of documents, the importer **108** starts the process by applying **410** the title pattern to a title of the document. If the importer **108** determines **420** that the title does not match the title pattern, it determines that the document should not be added to the set and repeats the process for the next document. Otherwise, the importer **108** applies **430** the contextual pattern to the content of the document. If the importer **108** determines **440** that there is no match for the contextual

pattern, it determines that the document should not be added to the set and repeats the process for the next document. If a title of the document matches the title pattern and the content of the document matches the contextual pattern, the importer **108** adds **450** the document to the set.

In one embodiment, the importer **108** only examines documents within the same domain as the source document when selecting **340** the set of documents. A domain is a group of computers and devices that share a common part of an Internet Protocol (IP) address that identifies (or locates) an organization (or entity) on the Internet. All documents from computers and devices in the group are said to be from the same domain. Documents from the same domain may be stored in one or multiple document hosts **102**. The set of documents can include the source document.

Limiting the set of documents to be within the same domain as the source document can be advantageous, especially when the size of the plurality of documents is large. The computational cost of comparing documents with the title pattern and the contextual pattern can become unmanageable as the number of documents grows. Also, because documents within the same domain as the source document are more likely to share the same style as the source document, they are more likely to match the title pattern and the contextual pattern than documents from other domains.

As described above, the importer **108** can identify **330** multiple contextual patterns. In one embodiment, the importer **108** adds **450** a document into the set of documents as long as the document matches the title pattern and any one of the multiple contextual patterns. In another embodiment, the importer **108** adds **450** a document into the set only when the document matches the title pattern and all of the identified contextual patterns.

Referring back to FIG. 3, for each document in the selected set, the importer **108** identifies **350** a name and one or more attribute-value pairs from the document. The importer **108** identifies **350** the name by applying the title pattern to the title of the document. The importer **108** identifies **350** one or more attribute-value pairs by applying the contextual pattern to the content of the document.

The identified name tends to be an object name for an object (or entity) described in the document. As described above, the source document contains facts of the source object, and therefore describes the source object. Further, the source object name can be extracted from a title of the source document using the title pattern. Because the title pattern also matches titles of the documents in the selected set, it follows that these documents tend to follow a consistent naming convention (or protocol or standard). For example, titles of web pages containing company profile in Yahoo! finance website follow a common pattern of “%COMPANY NAME% Company Profile—Yahoo! Finance”. Therefore, the name identified by applying the title pattern to a title of the document tends to be an object name for an object described in the document.

Similarly, the identified attribute-value pairs tend to be facts for the object described in the document. As described above, the source document contains facts of the source object matching the contextual pattern. Same as the source document, documents in the selected set match the title pattern and the contextual pattern. Therefore, these documents tend to share a style or format (e.g., displaying facts in a table or list). Because facts of the source object can be extracted from the source document using the contextual pattern, it follows that the attribute-value pairs identified by applying the contextual pattern to content of the document tend to be facts for the object described in the document.

11

The importer **108** can repeatedly apply the contextual pattern to content of the document to identify **350** (or extract or learn) attribute-value pairs. For example, if the contextual pattern identifies a table entry, the importer **108** may extract attribute-value pairs from that table entry and nearby table entries that match the contextual pattern. In one embodiment, the importer **108** extracts attribute-value pairs that are proximate to each other. By applying such a limitation, the importer **108** can be sure that the identified attribute-value pairs are likely from the same structure (e.g., the same table or list). Alternatively, the importer **108** can parse the document to identify **350** attribute-value pairs from all structures that match the contextual pattern.

In one embodiment, the importer **108** examines the quality of the identified attribute-value pairs by matching their attributes with a group of valid attributes (called “a white attribute list”). The identified contextual pattern may be general (e.g., a table row, a list item), and therefore the attribute-value pairs identified by applying the general contextual pattern may be over-inclusive. If the match fails to meet a certain threshold value (percentage or absolute value), the importer **108** determines that the identified attribute-value pairs are low in quality, and may disregard them entirely or only keep those that match. The white attribute list may include attributes identified in the source document and/or attributes collected from other sources (e.g., compiled by human editors). Similarly, the importer **108** may match the identified attribute-value pairs with a group of invalid attributes (called “a black attribute list”), and disregard them if the match exceeds a certain threshold value.

For each document in the selected set, the importer **108** associates **360** the attribute-value pairs identified in the document with an object with the identified name. In one embodiment, the importer **108** creates a new object, assigns the new object with the identified name, and associates **360** it with the identified attribute-value pairs. The importer **108** then passes the newly created object to a janitor **110** to incorporate the object into the fact repository **115**. In one embodiment, the janitor **110** searches in the fact repository **115** for an object associated with the identified name. If no such object exists, the janitor **110** adds the newly created object into the fact repository **115**. Otherwise, the janitor **110** merges the newly created object with the existing object in the fact repository **115**. If the janitor **110** determines that the existing object does not contain some of the attribute-value pairs identified in the document, it adds these attribute-value pairs into the fact repository **115** and associates **360** them with the existing object. If the existing object and the newly created object contain duplicated attribute-value pairs (e.g., same attribute and same value), the janitor **110** can optionally add the document as a source of this fact of the existing object. If the existing object and the newly created object contain inconsistent attribute-value pairs (e.g., same attribute but different value), the janitor **110** can determine a confidence value for each of the attribute-value pairs, and set the one with higher value as the value of the fact for the existing object. In other embodiments, the importer **108** only creates a new object if no object associated with the identified name is located in the fact repository **115**.

In one embodiment, the fact repository **115** can have objects associated with the same object name (or name fact value). For example, two objects can both have the name “New York,” one representing the New York City and the other representing a movie with the title “New York.” In order to locate the right object as described in the document, and not an object that happens to share the same object name, the janitor **110** searches in the fact repository **115** for an object

12

associated with the identified name and at least one identified attribute-value pair, and merges the newly created object with the object found in the manner detailed above.

It is noted that in alternative embodiments, the importer **108** (or the janitor **110**) may apply some normalization rules to the identified names to standardize its format before identifying corresponding objects. The importer **108** may also apply these normalization rules to the identified attribute-value pairs before associating **360** them with objects. Examples of the normalization rules include removal of punctuation, such as removing commas in a string, conversion of uppercase characters in a string to corresponding lowercase characters, such as from “America” to “america,” and stop word removal, such as removing stop words such as “the” and “of” from a string. For example, after applying the above normalization rules, an object name of “Charles Chaplin” becomes “charles chaplin”.

In one embodiment, the importer **108** (or the janitor **110**) compares the identified names with content of a black list. The black list contains texts that are not entity names (e.g., adjectives such as “beautiful,” and stop words such as “the”). Therefore, if there are matches between the identified name and the black list, the importer **108** determines that this name is not a valid entity name, and therefore does not create an object for it.

In one embodiment, the importer **108** (or the janitor **110**) compares the learned object names with content of a white list. The white list contains some entity names (e.g., “United Nation,” “John Lennon”). The white list can be compiled by human editors or extracted from reputable sources (e.g., the Encyclopedia Britannica Online). Therefore, if there are matches between the identified name and the white list, the importer **108** determines that the identified name is a proper object name.

The importer **108** can repeat the process described above for a different source object and/or a different source document. As a result, the importer **108** can learn objects and associated facts from the plurality of documents.

Example Process

The method **300** is further illustrated by an example shown in FIG. 5. As illustrated in FIG. 5, an object O1 has an attribute-value pair of “name-George Washington,” and an attribute-value pair of “Born-Feb. 22, 1732.” The source of both facts (or attribute-pairs) is “http://en.wikipedia.org/wiki/George_washington.”

FIG. 5 also shows a plurality of documents D1-D5. Documents D1-D5 are all from Wikipedia.org domain **501**. Document D1 (hereinafter called “D1”) is entitled “George Washington—Wikipedia, the free encyclopedia.” Content of D1 includes the following structured text:

Born	Feb. 22, 1732
Died	Dec. 14, 1799
Spouse	Martha Dandridge Custis Washington

Document D2 (“D2”) is entitled “Canada—Wikipedia, the free encyclopedia.” Content of D2 includes the following structured text:

Capital	Ottawa
Area	9,984,670 km ²
GDP	\$1.105 trillion

Document D3 (“D3”) is entitled “Help:Contents—Wikipedia, the free encyclopedia” and does not contain text struc-

tured as displayed above. Document D4 (“D4”) is entitled “Isaac Asimov—Wikipedia, the free encyclopedia.” Content of D4 includes the following structured text:

Born	Jan. 2, 1920
Died	Apr. 6, 1992

Document D5 (“D5”) is entitled “Fundraising—Wikimedia Foundation” and does not contain text structured as displayed above.

Content of documents D1-D5 can be organized in HTML structures such as tables and lists in order to be displayed as structured text as illustrated in FIG. 5. For example, the HTML source code of D1 may contain the following:

```
...
<title> George Washington—Wikipedia, the free encyclo-
pedia
</title>
...
<li> <b> Born </b> Feb. 22, 1732</li>
<li> <b> Died </b> Dec. 14, 1799</li>
<li> <b> Spouse </b> Martha Dandridge Custis Washing-
ton </li>
...
```

As shown above, content of D1 is organized as a list, and thus appears as structured text to viewers.

The importer 108 selects 310 object O1 as the source object. Because the source of both facts associated with O1 is “http://en.wikipedia.org/wiki/George_washington,” the URL of D1, the importer 108 selects 320 D1 as the source document.

The importer 108 identifies 330 a title pattern by comparing the title of the source document D1 (“George Washington—Wikipedia, the free encyclopedia”) with the object name of the source object O1 (“George Washington”). The importer 108 determines that the prefix section of the title pattern is empty and the suffix section is “—Wikipedia, the free encyclopedia”. Therefore, the identified title pattern is “%%&& %—Wikipedia, the free encyclopedia %”.

The importer 108 identifies 330 a contextual pattern by matching an attribute-value pair of the source object O1 (Born-Feb. 22, 1732) with the content of the source document D1 (“... Born Feb. 22, 1732 ...”). The importer 108 identifies 330 the following contextual pattern: “\$\$”.

The importer 108 selects 340 a set of documents based on the identified title pattern (“%%&& %—Wikipedia, the free encyclopedia %”) and the identified contextual pattern (“\$\$”). The importer 108 applies 410 the title pattern to titles of D1-D5 and determines 420 that titles of D1-D4 match with the title pattern. The importer 108 applies 430 the contextual pattern to contents of D1-D4, and determines 440 that the contents of D1, D2, and D4 match with the contextual pattern. As a result, the importer 108 adds 450 D1, D2, and D4 to the selected set of documents. It is noted that the selected set includes the source document D1.

The importer 108 applies the identified title pattern (“%%&& %—Wikipedia, the free encyclopedia %”) and contextual pattern (“\$\$”) to each document in the selected set (D1, D2, and D4) to identify 350 names and attribute-value pairs. For D1, the importer 108 identifies 350 the name “George Washington” and attribute value pairs “Born-Feb. 22, 1732,” “Died-Dec. 14, 1799,” and “Spouse-Martha Dandridge Custis Washington.” For D2, the importer 108 identifies the name “Canada” and attribute-value pairs “Capital-Ottawa,” “Area-9,984,670 km2,” and “GDP-\$1.105

trillion.” Similarly, for D4, the importer 108 identifies the name “Isaac Asimov” and attribute-value pairs “Born-Jan. 2, 1920” and “Died-Apr. 6, 1992.”

The importer 108 (or the janitor 110) identifies (or creates) objects with the identified names and associates 360 the identified attribute-value pairs with the objects. For D1, because O1 is already associated with the attribute-value pair “Born-Feb. 22, 1732,” the importer 108 (or the janitor 110) associates 360 the attribute-value pairs “Died-Dec. 14, 1799” and “Spouse-Martha Dandridge Custis Washington” with O1 and adds D1 as their source. For D2, the importer 108 (or the janitor 110) creates an object with object name “Canada,” associates 360 it with attribute-value pairs “Capital-Ottawa,” “Area-9,984,670 km2,” and “GDP-\$1.105 trillion,” and adds D2 as their source. Similarly, for D4, the importer 108 (or the janitor 110) creates an object with object name “Isaac Asimov,” associates 360 it with attribute-value pairs “Born-Jan. 2, 1920” and “Died-Apr. 6, 1992,” and adds D4 as their source.

Alternative Embodiments

In one embodiment, instead of selecting 320 a source document based on facts of the identified source object, the importer 108 can first select a source document and then select a source object based on the subject of the source document. For example, the importer 108 can first identify a title pattern of the selected source document, determine an object name based on the document title and the title pattern, and select a source object associated with the object name. The importer 108 can then identify 330 a contextual pattern as described above. If the importer 108 fails to select such a source object or fails to identify 330 a contextual pattern, it can repeat the process and select a different source document.

The importer 108 can identify a title pattern of the source document by analyzing the document title of the source document and anchor texts of documents linking to the source document. Detailed description and embodiments of identifying a title pattern for the source document can be found in U.S. patent application Ser. No. 11/394,610, entitled “Determining Document Subject by Using Title and Anchor Text of Related Documents,” filed on Mar. 31, 2006, the disclosure of which is hereby incorporated by reference in its entirety.

It is noted that the importer 108 can identify 330 more than one contextual pattern for the source document and the source object. For example, some of the attribute-value pairs may be in a list in the source document, and others may be in a table in the source document. In some embodiments, the importer 108 can select 340 a set of documents for each of the identified contextual patterns.

It should be appreciated that the process described above may be performed at another time for other purposes, such as learning images and other media data from semi-structured text and verifying the quality of the plurality of documents. The process may be performed as needed or at scheduled intervals. Furthermore, the process may be performed for other objects in the fact repository 115.

Reference in the specification to “one embodiment” or to “an embodiment” means that a particular feature, structure, or characteristic described in connection with the embodiments is included in at least one embodiment of the invention. The appearances of the phrase “in one embodiment” in various places in the specification are not necessarily all referring to the same embodiment.

Some portions of the above are presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory. These algorithmic descrip-

15

tions and representations are the means used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of steps (instructions) leading to a desired result. The steps are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical, magnetic or optical signals capable of being stored, transferred, combined, compared and otherwise manipulated. It is convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like. Furthermore, it is also convenient at times, to refer to certain arrangements of steps requiring physical manipulations of physical quantities as modules or code devices, without loss of generality.

It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. Unless specifically stated otherwise as apparent from the following discussion, it is appreciated that throughout the description, discussions utilizing terms such as "processing" or "computing" or "calculating" or "determining" or "displaying" or "determining" or the like, refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system memories or registers or other such information storage, transmission or display devices.

Certain aspects of the present invention include process steps and instructions described herein in the form of an algorithm. It should be noted that the process steps and instructions of the present invention can be embodied in software, firmware or hardware, and when embodied in software, can be downloaded to reside on and be operated from different platforms used by a variety of operating systems.

The present invention also relates to an apparatus for performing the operations herein. This apparatus may be specially constructed for the required purposes, or it may comprise a general-purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program may be stored in a computer readable storage medium, such as, but is not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, magnetic-optical disks, read-only memories (ROMs), random access memories (RAMs), EPROMs, EEPROMs, magnetic or optical cards, application specific integrated circuits (ASICs), or any type of media suitable for storing electronic instructions, and each coupled to a computer system bus. Furthermore, the computers referred to in the specification may include a single processor or may be architectures employing multiple processor designs for increased computing capability.

The algorithms and displays presented herein are not inherently related to any particular computer or other apparatus. Various general-purpose systems may also be used with programs in accordance with the teachings herein, or it may prove convenient to construct more specialized apparatus to perform the required method steps. The required structure for a variety of these systems will appear from the description below. In addition, the present invention is not described with reference to any particular programming language. It will be appreciated that a variety of programming languages may be used to implement the teachings of the present invention as described herein, and any references below to specific languages are provided for disclosure of enablement and best mode of the present invention.

16

While the invention has been particularly shown and described with reference to a preferred embodiment and several alternate embodiments, it will be understood by persons skilled in the relevant art that various changes in form and details can be made therein without departing from the spirit and scope of the invention.

Finally, it should be noted that the language used in the specification has been principally selected for readability and instructional purposes, and may not have been selected to delineate or circumscribe the inventive subject matter. Accordingly, the disclosure of the present invention is intended to be illustrative, but not limiting, of the scope of the invention, which is set forth in the following claims.

What is claimed is:

1. A method for learning objects and facts from documents, comprising:

on a computer system having one or more processors and memory storing one or more programs for execution by the one or more processors, without user intervention:

selecting a source object from a plurality of objects stored in a fact repository, the source object having an object name and an attribute-value pair;

selecting a source document from a plurality of documents based on a determination that

a title of the source document includes the object name of the source object; and

content of the source document includes an attribute and a value related to the attribute-value pair of the source object;

identifying a title pattern for the title of the source document based on a first syntax of (i) the title of the source document, and (ii) the object name of the source object within the title of the source document;

identifying a contextual pattern for the content of the source document based on a second syntax of (i) the attribute and (ii) the value related to the attribute-value pair of the source object, the contextual pattern being a structural pattern in which one or more attribute-value pairs, including the attribute-value pair of the source object, are presented;

selecting a second document from the plurality of documents based on a determination that (i) the contextual pattern is found in the second document and (ii) a title of the second document matches the title pattern;

identifying a new object name and a new attribute-value pair from the second document by applying the title pattern and the contextual pattern identified from the source document to the second document; and

storing into the fact repository the new attribute-value pair and a new object having the new object name.

2. The method of claim 1, wherein the first syntax specifies a first sequence of symbols and non-symbol characters that identifies the object name within the title of the source document.

3. The method of claim 1, wherein selecting the source document from the plurality of documents further comprises selecting a source of the attribute-value pair of the source object as the source document.

4. The method of claim 1, wherein selecting the source document from the plurality of documents further comprises: obtaining a search result including documents containing the object name and the attribute-value pair of the source object; and selecting the source document from the search result.

5. The method of claim 1, wherein selecting the second document further comprises selecting a set of documents from the plurality of documents, each document in the

17

selected set matching the title pattern and the contextual pattern, the method further comprising:

repeating the operation of identifying the new object name and the new attribute-value pair for each of the documents in the selected set.

6. The method of claim 5, wherein the documents in the selected set are from a same domain as the source document.

7. The method of claim 5, wherein the documents in the selected set include the source document.

8. The method of claim 1, wherein identifying the new object name and the new attribute-value pair further comprises:

matching an attribute of the new attribute-value pair with a white attribute list;

responsive to a match, determining that the new attribute-value pair is a valid attribute-value pair; and

responsive to no match, determining that the new attribute-value pair is an invalid attribute-value pair.

9. The method of claim 1, wherein storing the new attribute-value pair and the new object having the new object name further comprises:

identifying the new object associated with the new object name; and

associating the new attribute-value pair with the new object.

10. The method of claim 9, wherein identifying the new object associated with the new object name comprises:

responsive to an existing object comprising the new object name, identifying the existing object as the new object; and

responsive to no existing object comprising the new object name:

creating an object,

assigning the new object name to the object, and

identifying the object as the new object.

11. The method of claim 1, wherein the plurality of documents comprise a web page.

12. The method of claim 1, wherein the source document and the second document share (i) a same title pattern for their respective title, and (ii) a same contextual pattern for their respective content.

13. The method of claim 1, wherein title of the source document includes (i) a prefix section, (ii) the object name of the source object, and (iii) a suffix section; and

identifying the title pattern for the title of the source document based on structural arrangement the object name of the source object includes:

identifying the title pattern for the title of the source document based on structural arrangement of (i) the prefix section, (ii) the object name of the source object, and (iii) the suffix section within the title of the source document.

14. The method of claim 1, wherein identifying the title pattern for the title of the source document based on structural arrangement the object name of the source object includes:

identifying, within the title of the source document, a substring corresponding to the object name of the source object; and

determining, position of the substring relative to one or more other strings in the title of the source document.

15. A system for learning objects and facts from document, the system comprising:

a processor;

memory; and

one or more programs for execution by the processor, wherein the one or more programs including instructions for, without user intervention:

18

selecting a source object from a plurality of objects stored in a fact repository, the source object having an object name and an attribute-value pair;

selecting a source document from a plurality of documents based on a determination that

a title of the source document includes the object name of the source object; and

content of the source document includes an attribute and a value related to the attribute-value pair of the source object;

identifying a title pattern for the title of the source document based on a first syntax of (i) the title of the source document, and (ii) the object name of the source object within the title of the source document;

identifying a contextual pattern for the content of the source document based on a second syntax of (i) the attribute and (ii) the value related to the attribute-value pair of the source object, the contextual pattern being a structural pattern in which one or more attribute-value pairs, including the attribute-value pair of the source object, are presented;

selecting a second document from the plurality of documents based on a determination that (i) the contextual pattern is found in the second document and (ii) a title of the second document matches the title pattern;

identifying a new object name and a new attribute-value pair from the second document by applying the title pattern and the contextual pattern identified from the source document to the second document; and

storing into the fact repository the new attribute-value pair and a new object having the new object name.

16. The system of claim 15, wherein the first syntax specifies a first sequence of symbols and non-symbol characters that identifies the object name within the title of the source document.

17. The system of claim 15, wherein selecting the second document further comprises selecting a set of documents from the plurality of documents, each document in the selected set matching the title pattern and the contextual pattern, the method further comprising:

repeating the operation of identifying the new object name and the new attribute-value pair for each of the documents in the selected set.

18. A non-transitory computer readable storage medium storing one or more programs configured for execution by a computer, the one or more programs comprising instructions for, without user intervention:

selecting a source object from a plurality of objects stored in a fact repository, the source object having an object name and an attribute-value pair;

selecting a source document from a plurality of documents based on a determination that

a title of the source document includes the object name of the source object; and

content of the source document includes an attribute and a value related to the attribute-value pair of the source object;

identifying a title pattern for the title of the source document based on a first syntax of (i) the title of the source document, and (ii) the object name of the source object within the title of the source document;

identifying a contextual pattern for the content of the source document based on a second syntax of (i) the attribute and (ii) the value related to the attribute-value pair of the source object, the contextual pattern being a

structural pattern in which one or more attribute-value pairs, including the attribute-value pair of the source object, are presented;

selecting a second document from the plurality of documents based on a determination that (i) the contextual 5 pattern is found in the second document and (ii) a title of the second document matches the title pattern;

identifying a new object name and a new attribute-value pair from the second document by applying the title pattern and the contextual pattern identified from the 10 source document to the second document; and

storing into the fact repository the new attribute-value pair and a new object having the new object name.

19. The non-transitory computer readable storage medium of claim **18**, wherein the first syntax specifies a first sequence 15 of symbols and non-symbol characters that identifies the object name within the title of the source document.

20. The non-transitory computer readable storage medium of claim **18**, wherein the instructions for selecting the second document further comprises instructions for selecting a set of 20 documents from the plurality of documents, each document in the selected set matching the title pattern and the contextual pattern, the one or more computer programs further comprising:

instructions for repeating the operation of identifying the 25 new object name and the new attribute-value pair for each of the documents in the selected set.

* * * * *